

# User Level Failure Mitigation in MPI

Wesley Bland

MPI Forum Fault Tolerance Working Group

Resilience Workshop

8/28/12



# Motivation

- Failures are common enough that we want to handle them but rare enough that failure free performance is most important right now.
- Want plan for managing failures to be
  - Lightweight
  - Efficient
  - Extensible



“...faults will be continuous and across all parts of the hardware and software layers, which will require new programming paradigms.”

# How to Minimize Cost

# How to Minimize Cost

Collective operations return uniformly

```
for (i = 0; i < n; i++) {  
  
    value = compute();  
  
    rc = MPI_Bcast(comm, value);  
  
    if (rc!=MPI_SUCCESS)  
  
        recover();  
  
}
```

# How to Minimize Cost

Collective operations return uniformly

```
for (i = 0; i < n; i++) {  
    value = compute();  
    rc = MPI_Bcast(comm, value);  
    if (rc!=MPI_SUCCESS)  
        recover();  
}
```

Bcast agrees on return code

$O(n^2)$

# How to Minimize Cost

Collective operations return uniformly

```
for (i = 0; i < n; i++) {  
  
    value = compute();  
  
    rc = MPI_Bcast(comm, value);  
  
    if (rc!=MPI_SUCCESS)  
  
        recover();  
  
}
```

Operations notify of any process failure

```
while(morework) {  
  
    rc = MPI_Recv(myrank-1);  
  
    if (rc!=MPI_SUCCESS) recover();  
  
    rc = MPI_Send(myrank+1);  
  
    if (rc!=MPI_SUCCESS) recover();  
}
```

Bcast agrees on return code

$O(n^2)$

# How to Minimize Cost

myrank-2

Collective operations return uniformly

```
for (i = 0; i < n; i++) {  
  
    value = compute();  
  
    rc = MPI_Bcast(comm, value);  
  
    if (rc!=MPI_SUCCESS)  
  
        recover();  
  
}
```

Bcast agrees on return code

$O(n^2)$

Operations notify of any process failure

```
while(morework) {  
  
    rc = MPI_Recv(myrank-1);  
  
    if (rc!=MPI_SUCCESS) recover();  
  
    rc = MPI_Send(myrank+1);  
  
    if (rc!=MPI_SUCCESS) recover();  
  
}
```

Failures must be reported  
during unrelated operations

# Two Guiding Principles

- A correct MPI program must not deadlock because of process failures.
- An MPI call that does not involve a failed process will complete normally unless we specifically intervene.

# Simple Master/Worker

- FT Method
  - Check Return Codes
- New MPI Constructs
- MPI\_ERR\_PROC\_FAILED

```
MPI_Irecv( /* all workers */ );  
  
while (work_available &&  
       active_workers > 0  
     )  
  {  
    rc = MPI_Waitany( &worker );  
  
    if (MPI_ERR_PROC_FAILED == rc)  
      {  
        active_workers--;  
  
        /* Recover Lost Work */  
      }  
    else  
      {  
        /* Send new work to worker */  
  
        MPI_Irecv( worker );  
      }  
  }
```

# Failure Acknowledgement

- `MPI_COMM_FAILURE_ACK( comm )`
- `MPI_COMM_FAILURE_GET_ACKED( comm, &group )`
- Discover the group of process which are known to have failed in **comm**
- Unmatched `MPI_ANY_SOURCE` receptions will not return failure for the processes in **group**
- **Group** is updated during the next call to ACK

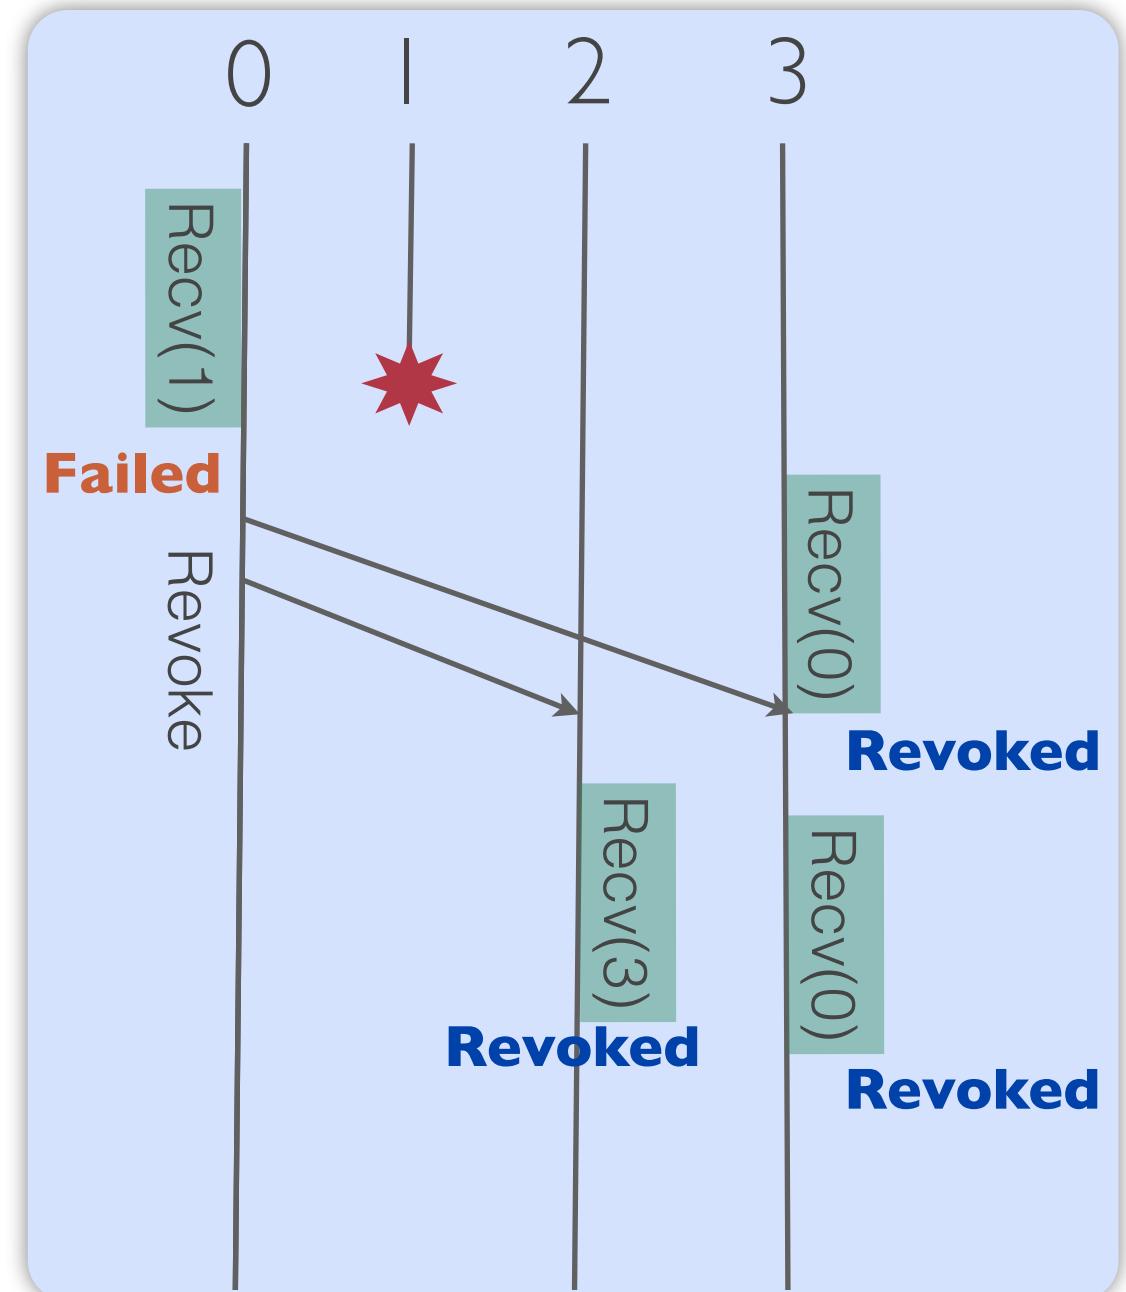
# Slightly harder...

- FT Method
  - Check return codes (MPI\_ERR\_PENDING)
  - Acknowledge known failures to continue with existing processes
- New MPI Constructs
  - MPI\_Comm\_failure\_ack
  - MPI\_Comm\_failure\_get\_acked

```
MPI_Irecv(MPI_ANY_SOURCE);  
  
while (work_available &&  
       active_workers > 0)  
    rc = MPI_Wait();  
  
    if (MPI_ERR_PENDING == rc  
        || MPI_ERR_PROC_FAILED == rc)  
        MPI_Comm_failure_ack(comm);  
  
        MPI_Comm_failure_get_acked(comm,  
                                    &group);  
  
        MPI_Group_size(group, &gsize);  
  
        active_workers -= /* Size difference */  
    else  
  
        MPI_Irecv(MPI_ANY_SOURCE);
```

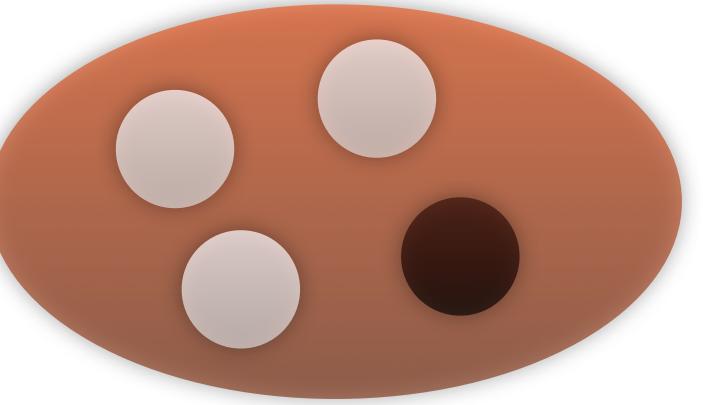
# MPI\_COMM\_REVOKE(comm)

- Non-collective operation
- All non-local operations complete with error
- New operations will always return an error
- Communicator is revoked when the local process calls REVOKE or returns MPI\_ERR\_REVOKED because another process called it



# **MPI\_COMM\_SHRINK(comm, comm2)**

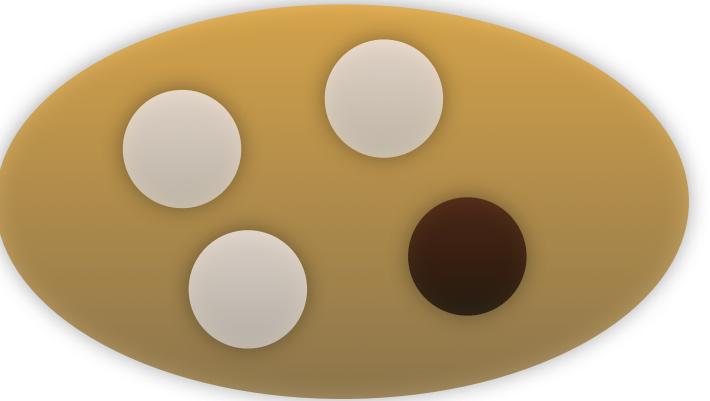
- Creates a new communicator from a revoked one
- Collective over non-failed processes
- Will not return an error due to process failure
- Results in two communicators
  - Dead & Shrunken



# **MPI\_COMM\_SHRINK(comm, comm2)**

- Creates a new communicator from a revoked one
- Collective over non-failed processes
- Will not return an error due to process failure
- Results in two communicators
  - Dead & Shrunken

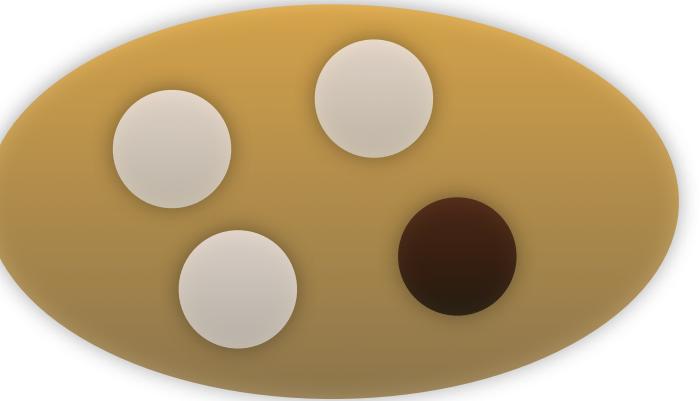
**MPI\_Revoke**



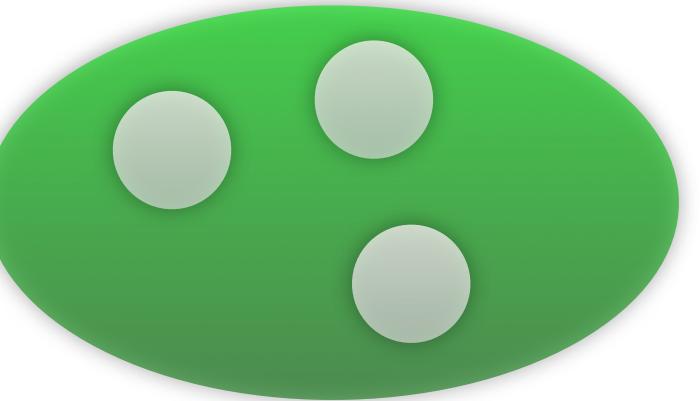
# **MPI\_COMM\_SHRINK(comm, comm2)**

- Creates a new communicator from a revoked one
- Collective over non-failed processes
- Will not return an error due to process failure
- Results in two communicators
  - Dead & Shrunken

**MPI\_Revoke**

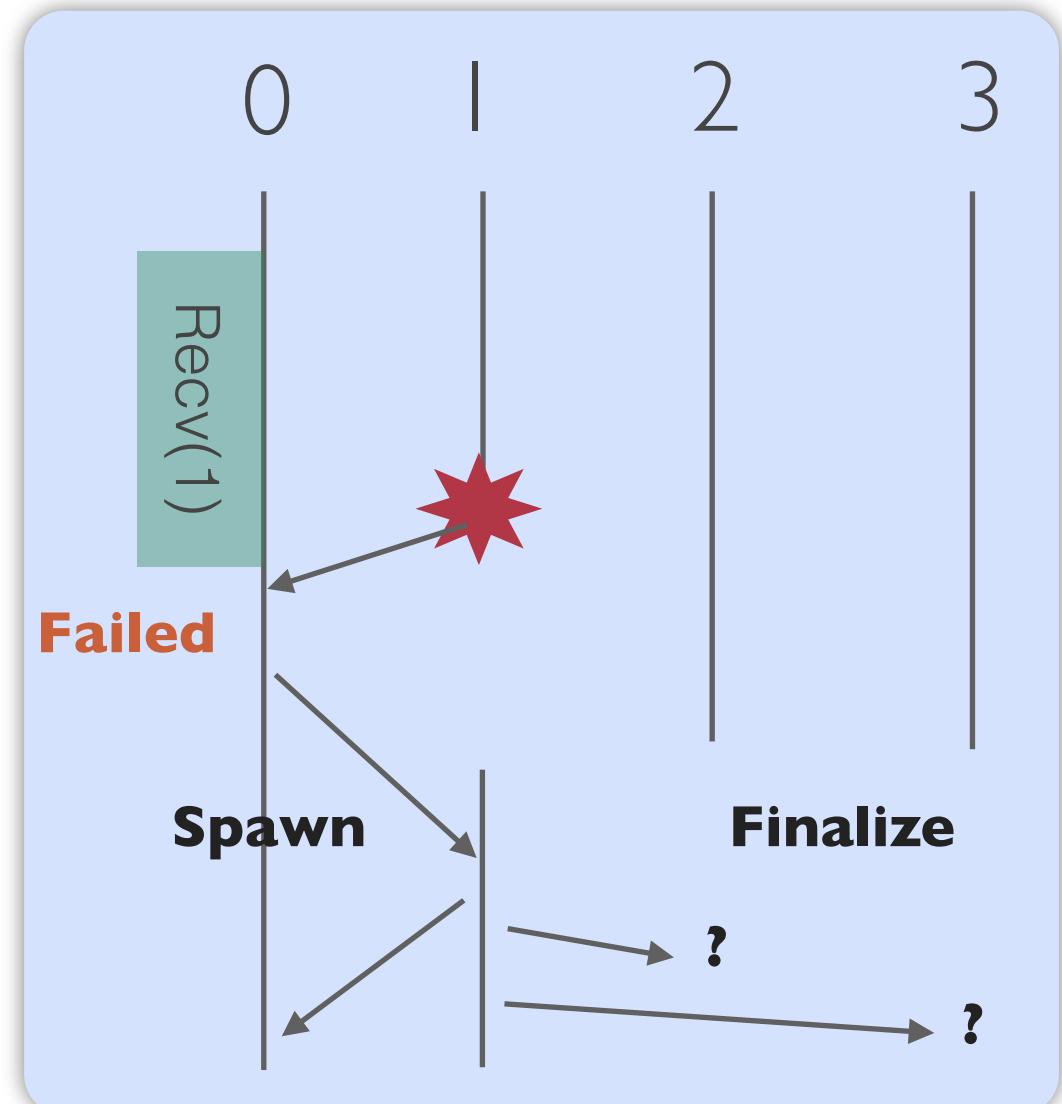


**MPI\_Shrink**



# **MPI\_COMM\_AGREE(comm, flag)**

- All living processes perform logical AND over ***flag***
- Failed processes do not contribute
- Will not return error due to process failure
- Revoked communicators uniformly return MPI\_ERR\_REVOKED



# Error Codes

- **MPI\_ERR\_PROC\_FAILED**
  - Operation cannot complete because involved process is dead
- **MPI\_ERR\_REVOKED**
  - Communication object is revoked and will not work again
- **MPI\_ERR\_PENDING**
  - Non-blocking completion **may** not complete because involved process is dead

# Iterative Refinement

```
while( gnorm > epsilon ) {  
    rc = MPI_Allreduce( &lnorm, &gnorm, 1, MPI_DOUBLE, MPI_MAX,  
    comm);  
    if( (MPI_ERR_PROC_FAILED == rc) ||  
        (MPI_ERR_COMM_REVOKED == rc) ||  
        (gnorm <= epsilon) ) {  
  
        allsucceeded = (rc == MPI_SUCCESS);  
        MPI_Comm_agree(comm, &allsucceeded);  
        if( !allsucceeded ) {  
            MPI_Comm_revocate(comm);  
            MPI_Comm_shrink(comm, &comm2);  
            MPI_Comm_free(comm);  
            comm = comm2;  
  
            gnorm = epsilon + 1.0;
```

```
        rc = MPI_Allreduce( &lnorm, &gnorm, 1, MPI_DOUBLE, MPI_MAX,  
comm);  
  
    if( (MPI_ERR_PROC_FAILED == rc) ||  
        (MPI_ERR_COMM_REVOKED == rc) ||  
        (gnorm <= epsilon) ) {  
  
        allsucceeded = (rc == MPI_SUCCESS);  
        MPI_Comm_agree(comm, &allsucceeded);  
        if( !allsucceeded ) {  
            MPI_Comm_revoke(comm);  
            MPI_Comm_shrink(comm, &comm2);  
            MPI_Comm_free(comm);  
            comm = comm2;  
            gnorm = epsilon + 1.0;  
        } } }
```

# This seems weak...

- Stronger consistency models can sit on top of MPI inside a library
- Most flexibility for the user
- No “silver bullet”

Application

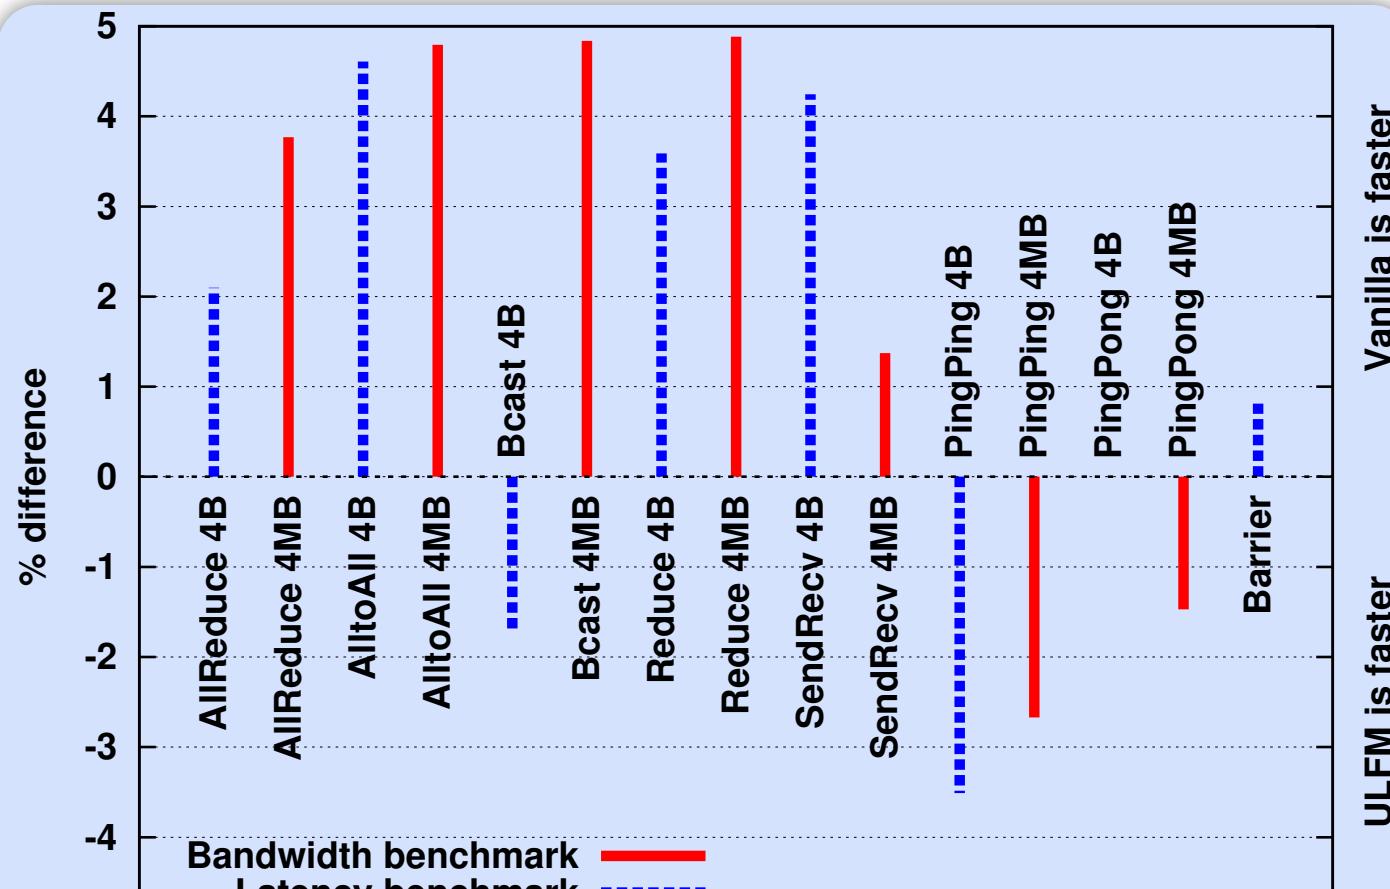
Transactions

Uniform  
Collectives

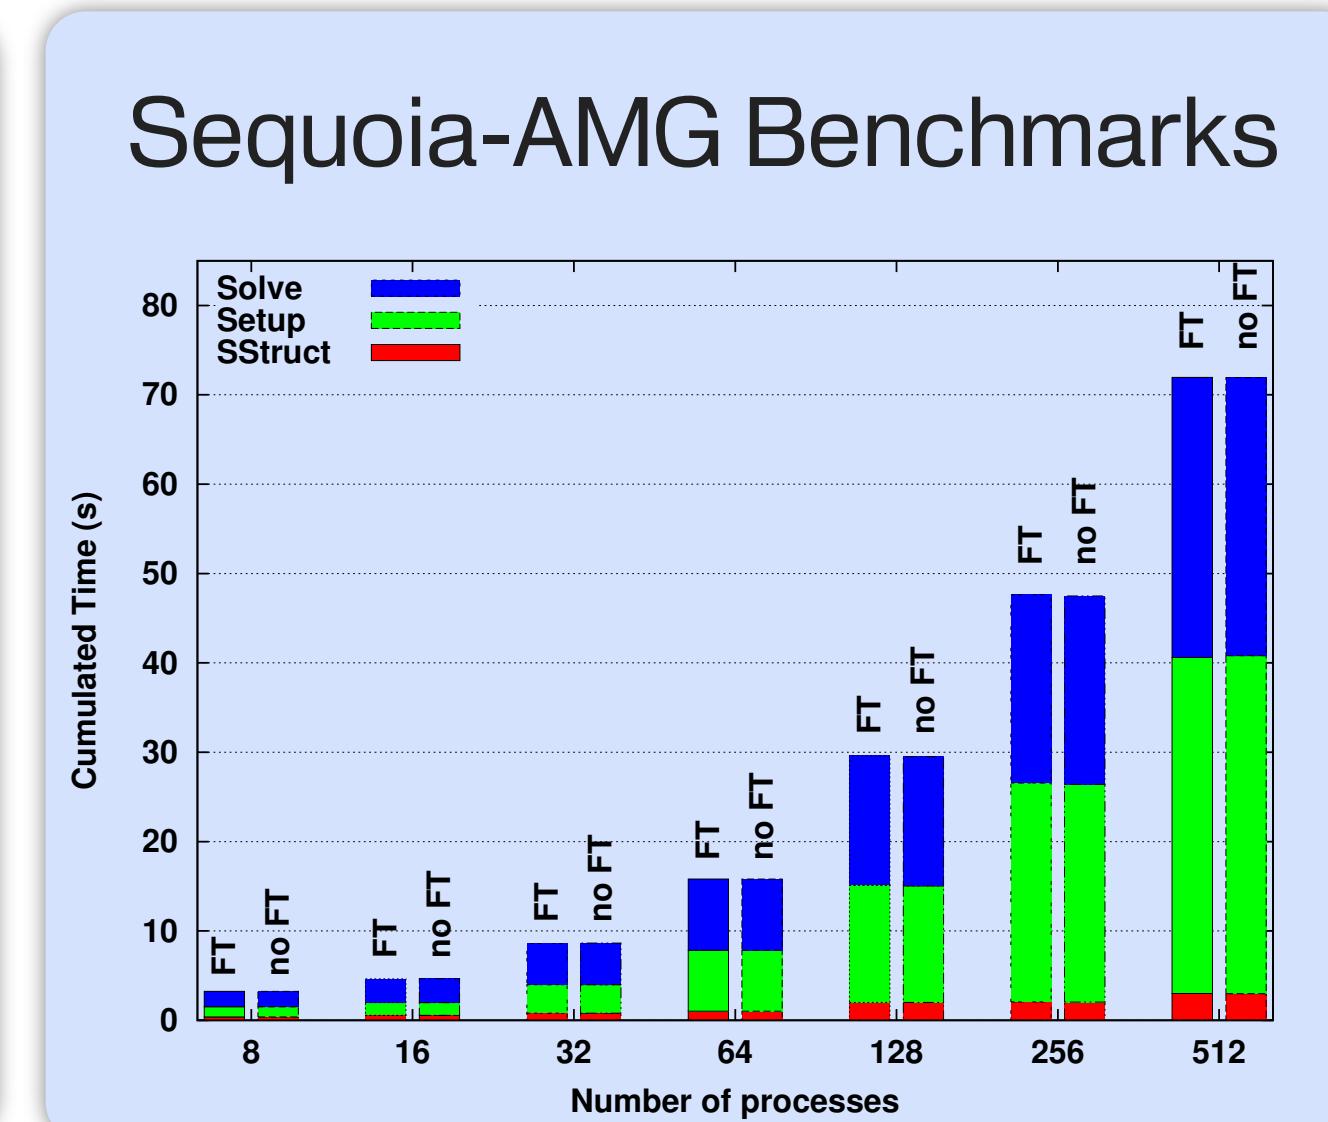
Others

MPI

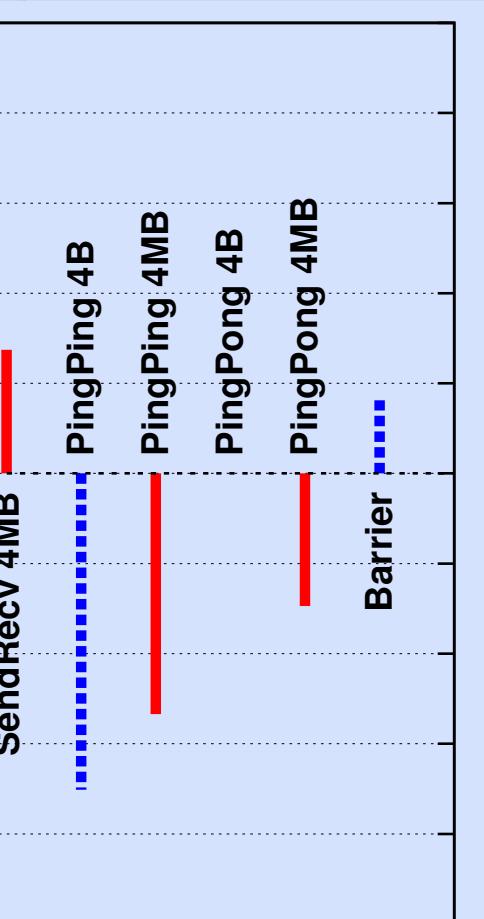
# ULFM Results



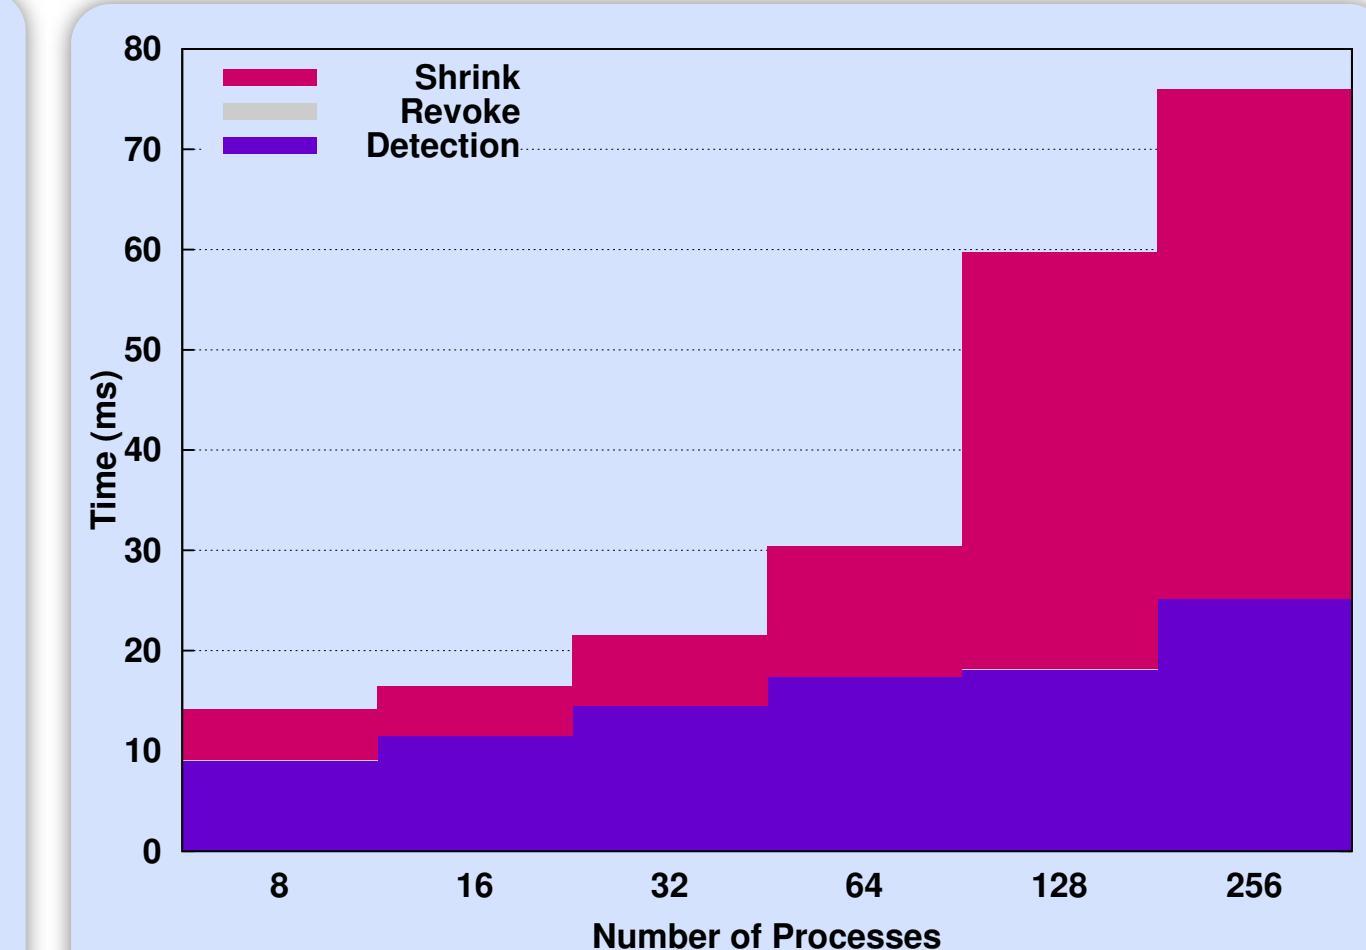
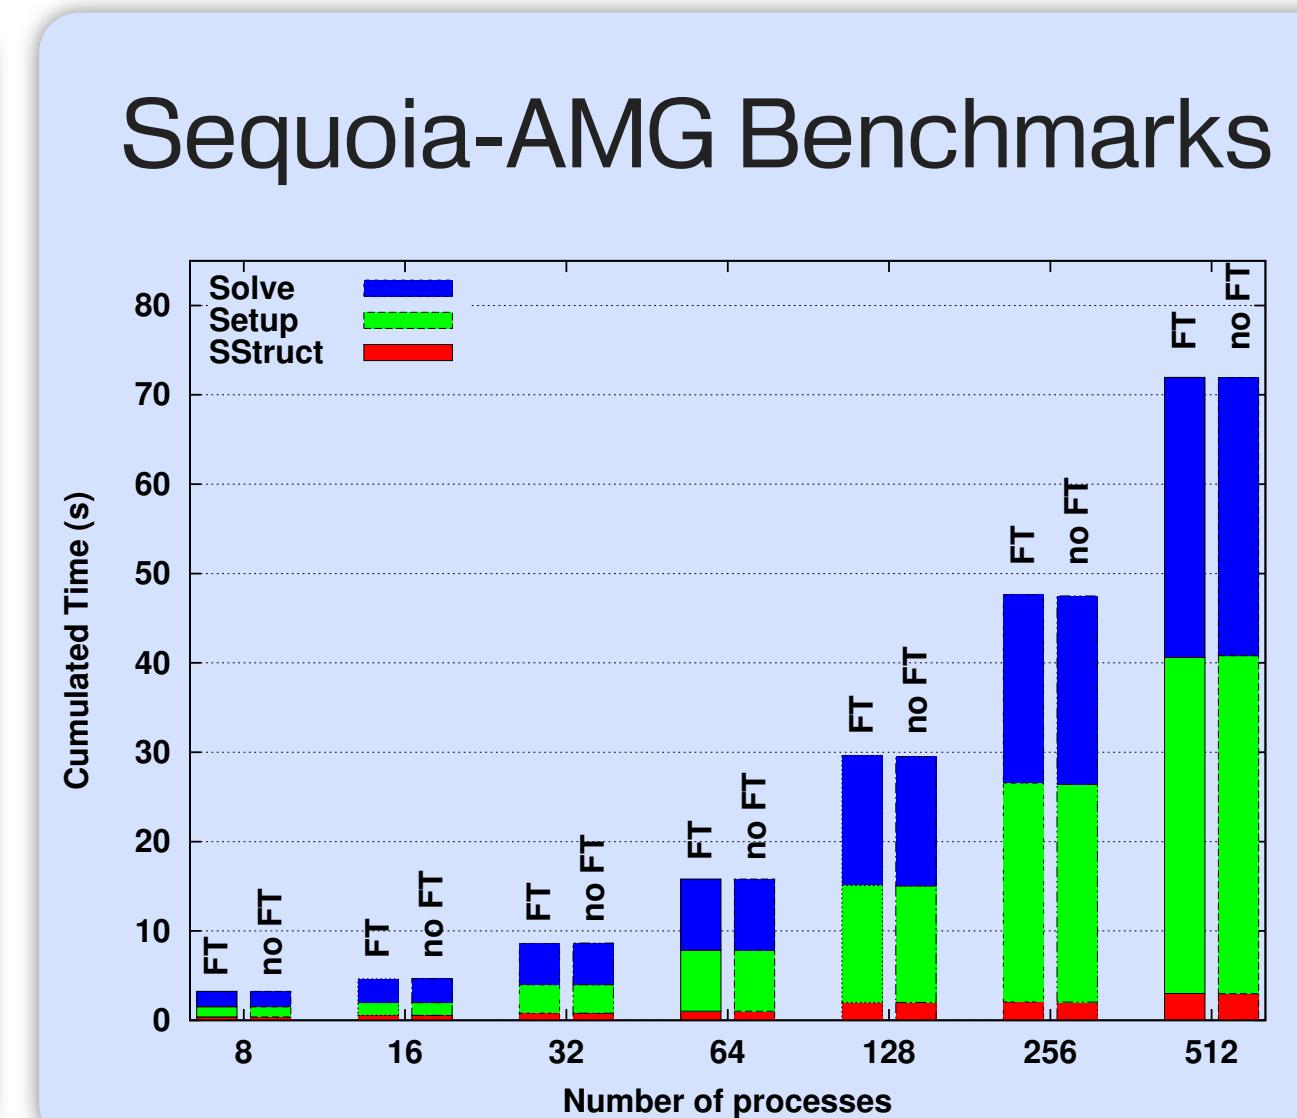
Intel MPI Benchmarks



# ULFM Results



benchmarks



Synthetic Benchmark

# How do we use this?

- 3 R's
- Library Composition
- Real Applications

# 3 R's of Application Recovery

- **Revoke**
  - Revoke all communicators involved with failed processor(s)
- **Repair**
  - Repair (or replace) the revoked communicators and recover the lost data
- **Resume**
  - Resume execution from the interrupted point (or the most recent recovery data)

# Library Composition

MAIN:

```
main( ) { }  
repair( ) { }
```

```
MPI_Comm my_comm_world;
```

LIBRARY 1:

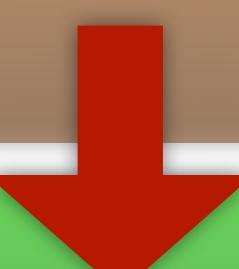
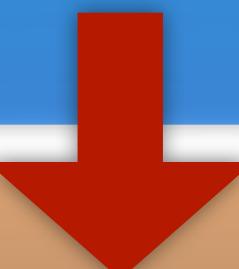
```
function( ) { }  
repair( ) { }
```

```
MPI_Comm lib1_comm;
```

LIBRARY 2:

```
function( ) { }  
repair( ) { }
```

```
MPI_Comm lib2_comm;
```



# Library Composition

MAIN:

```
main() { }  
repair() { }
```

~~MPI\_Comm my\_comm\_world;~~

LIBRARY 1:

```
function() { }  
repair() { }
```

~~MPI\_Comm lib1\_comm;~~

LIBRARY 2:

```
function() { }  
repair() { }
```

~~MPI\_Comm lib2\_comm;~~

# Library Composition

MAIN:

```
main() { }  
repair() { }
```

```
MPI_Comm my_comm_world;
```

LIBRARY 1:

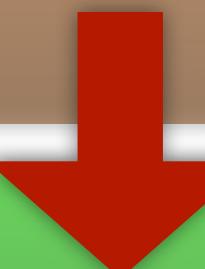
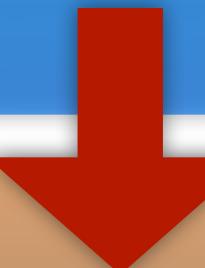
```
function() { }  
repair() { }
```

```
MPI_Comm lib1_comm;
```

LIBRARY 2:

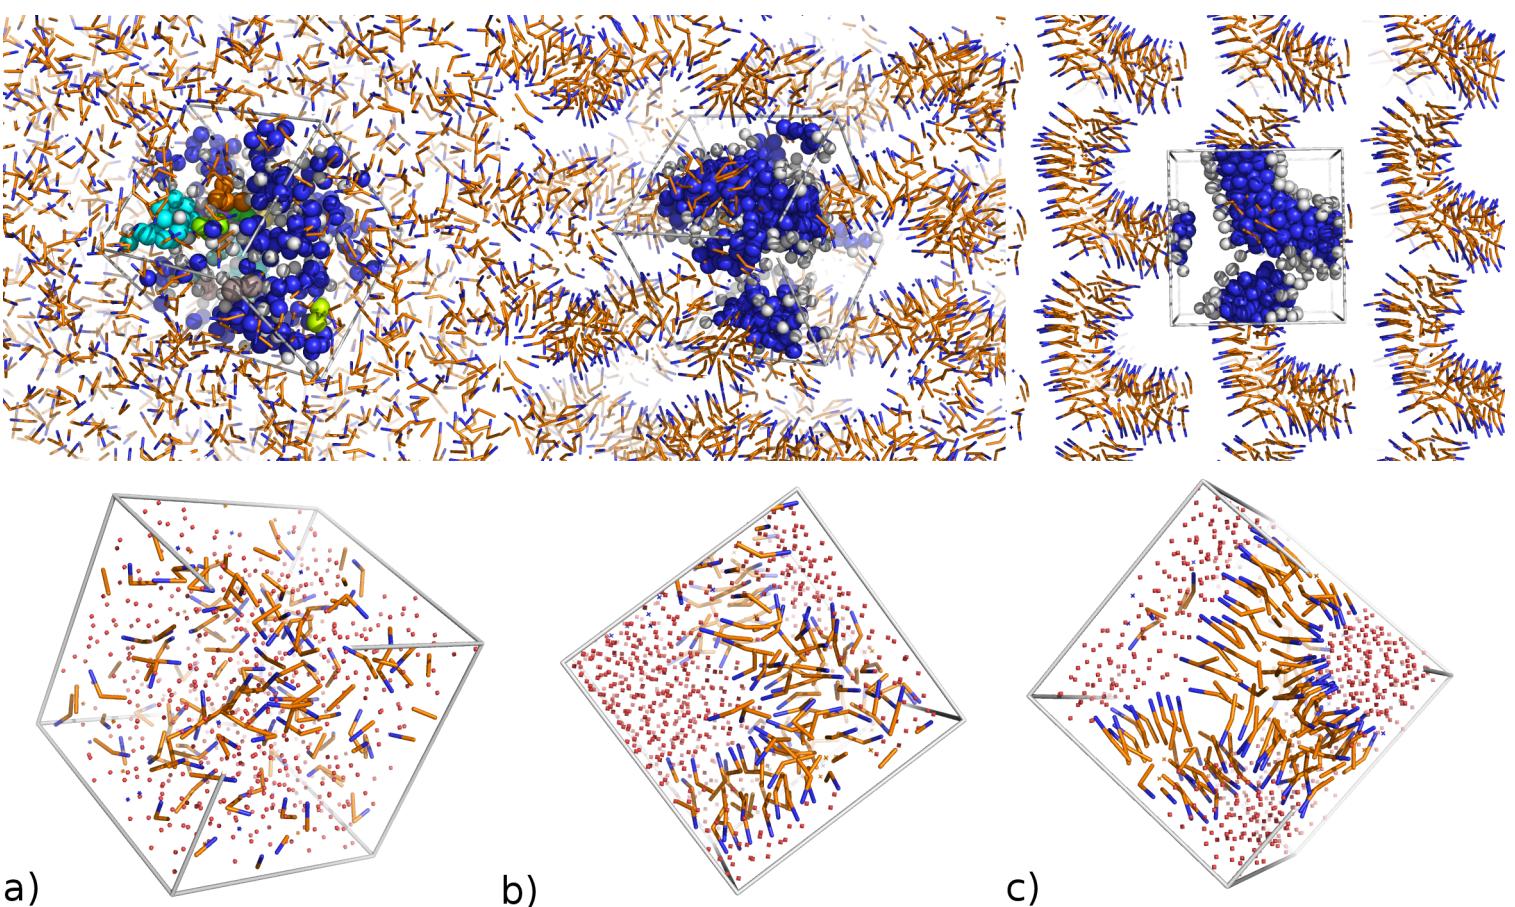
```
function() { }  
repair() { }
```

```
MPI_Comm lib2_comm;
```



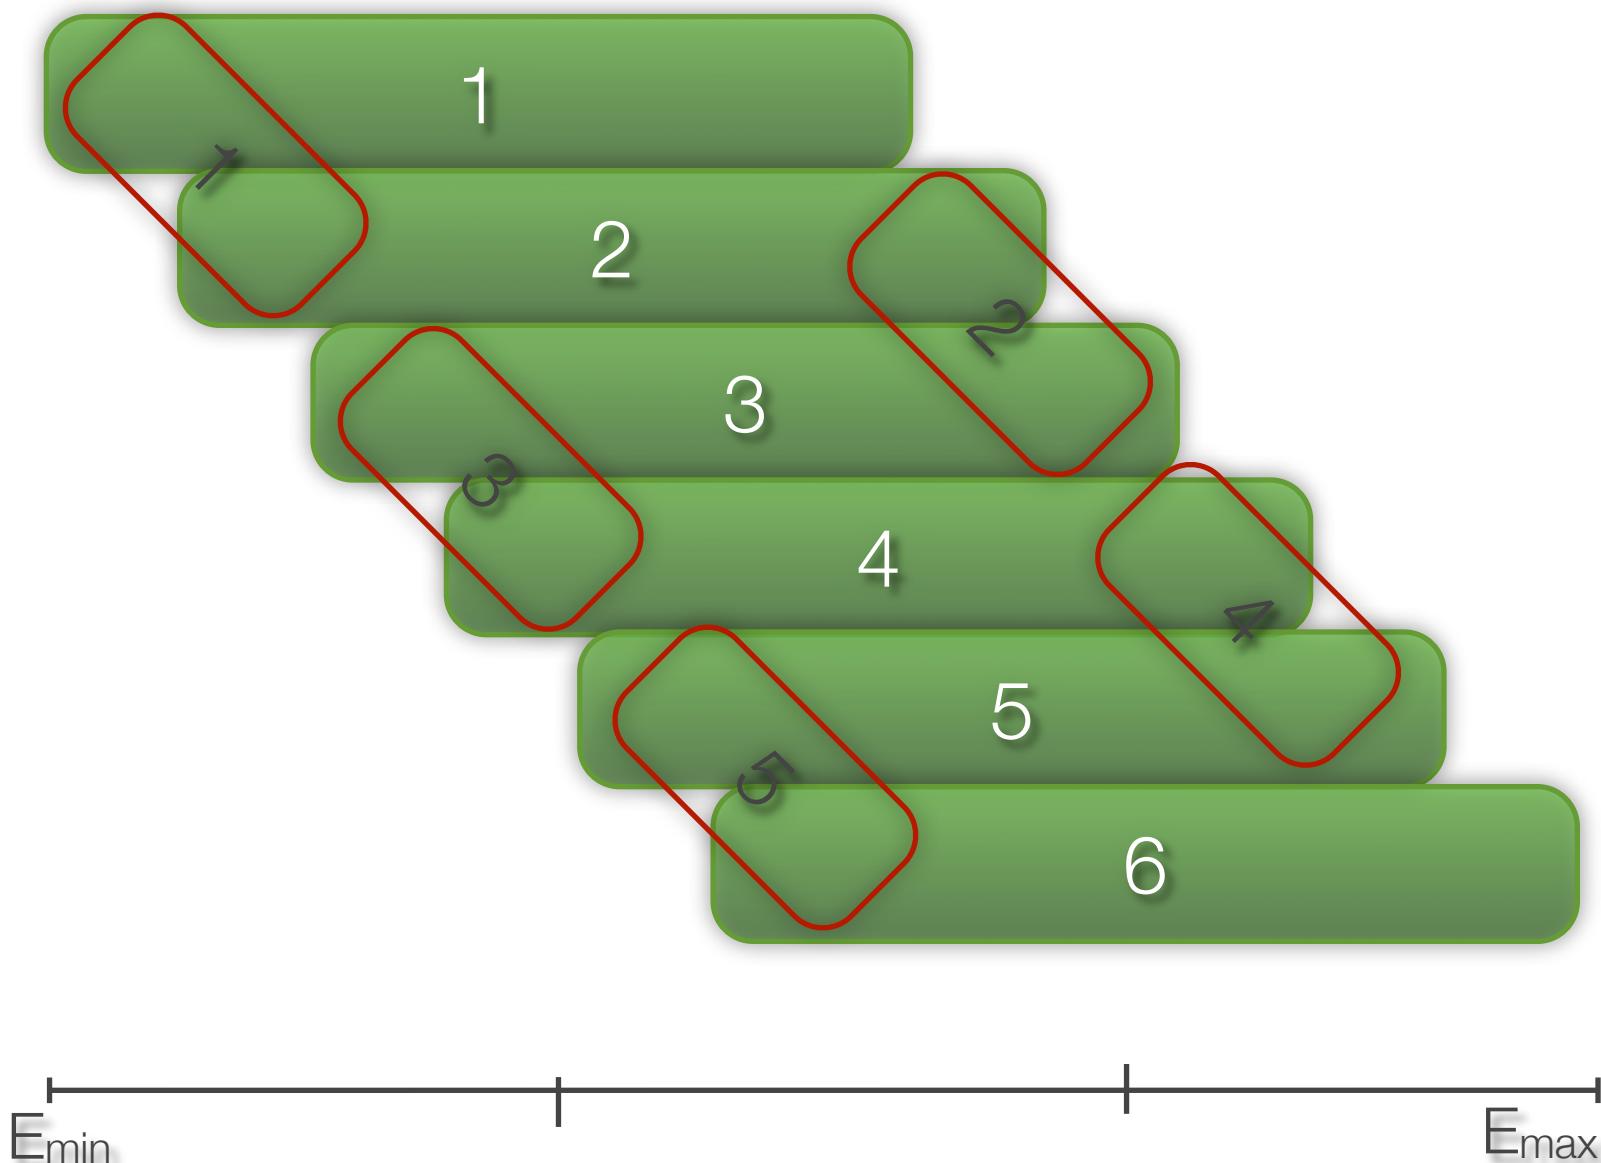
# Computation Chemistry Application

- Developed at the University of Georgia
- Simulating molecular interaction in skin cells
- Uses two types of communicators
  - Communication within an energy window
  - Communication between energy windows



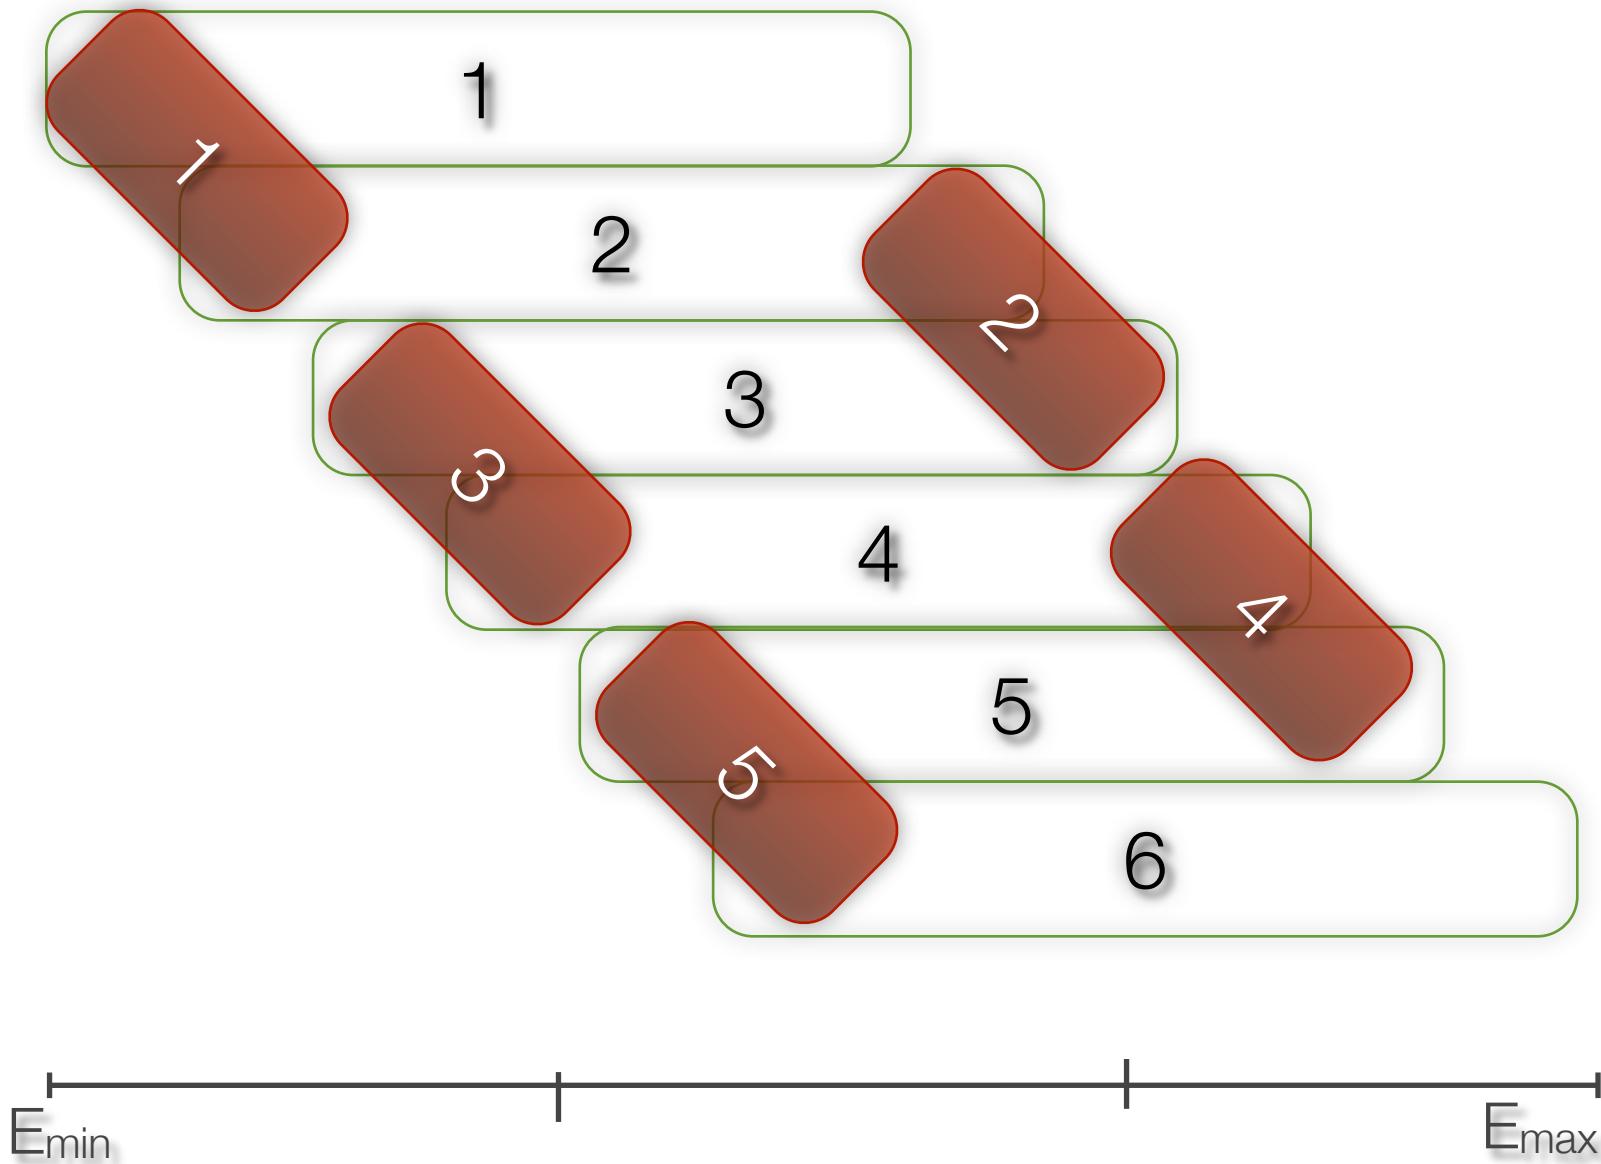
# Application Model

- 'Green' Communicator
  - Includes all processes in own energy window
  - Communication:
    - 2x MPI\_Allreduce

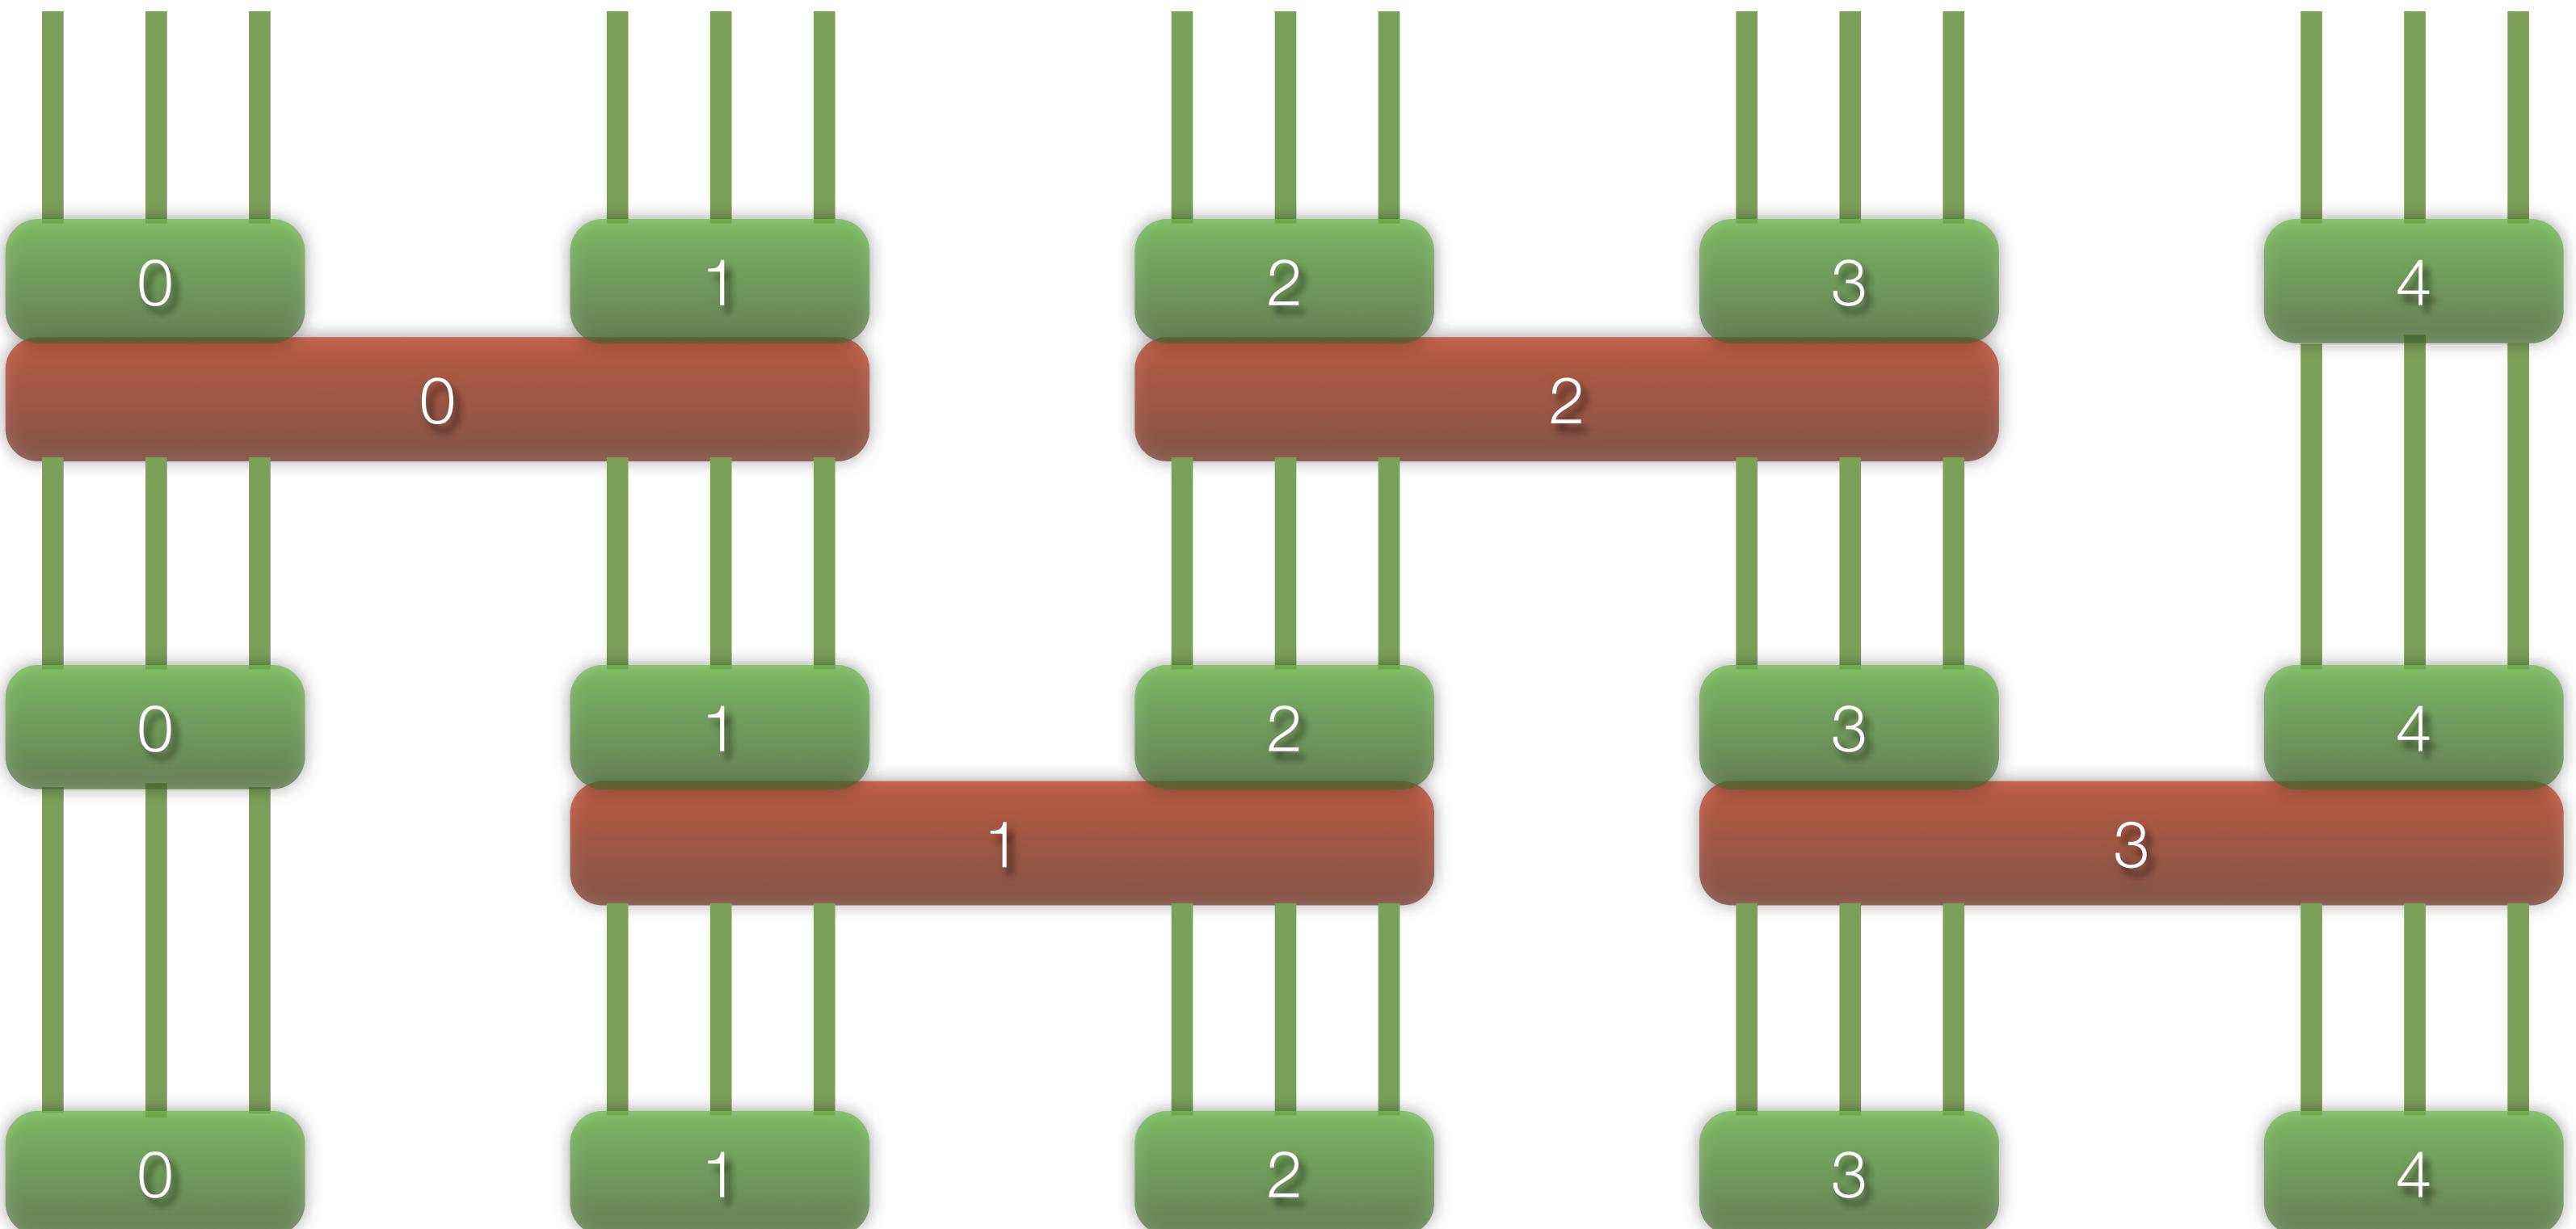


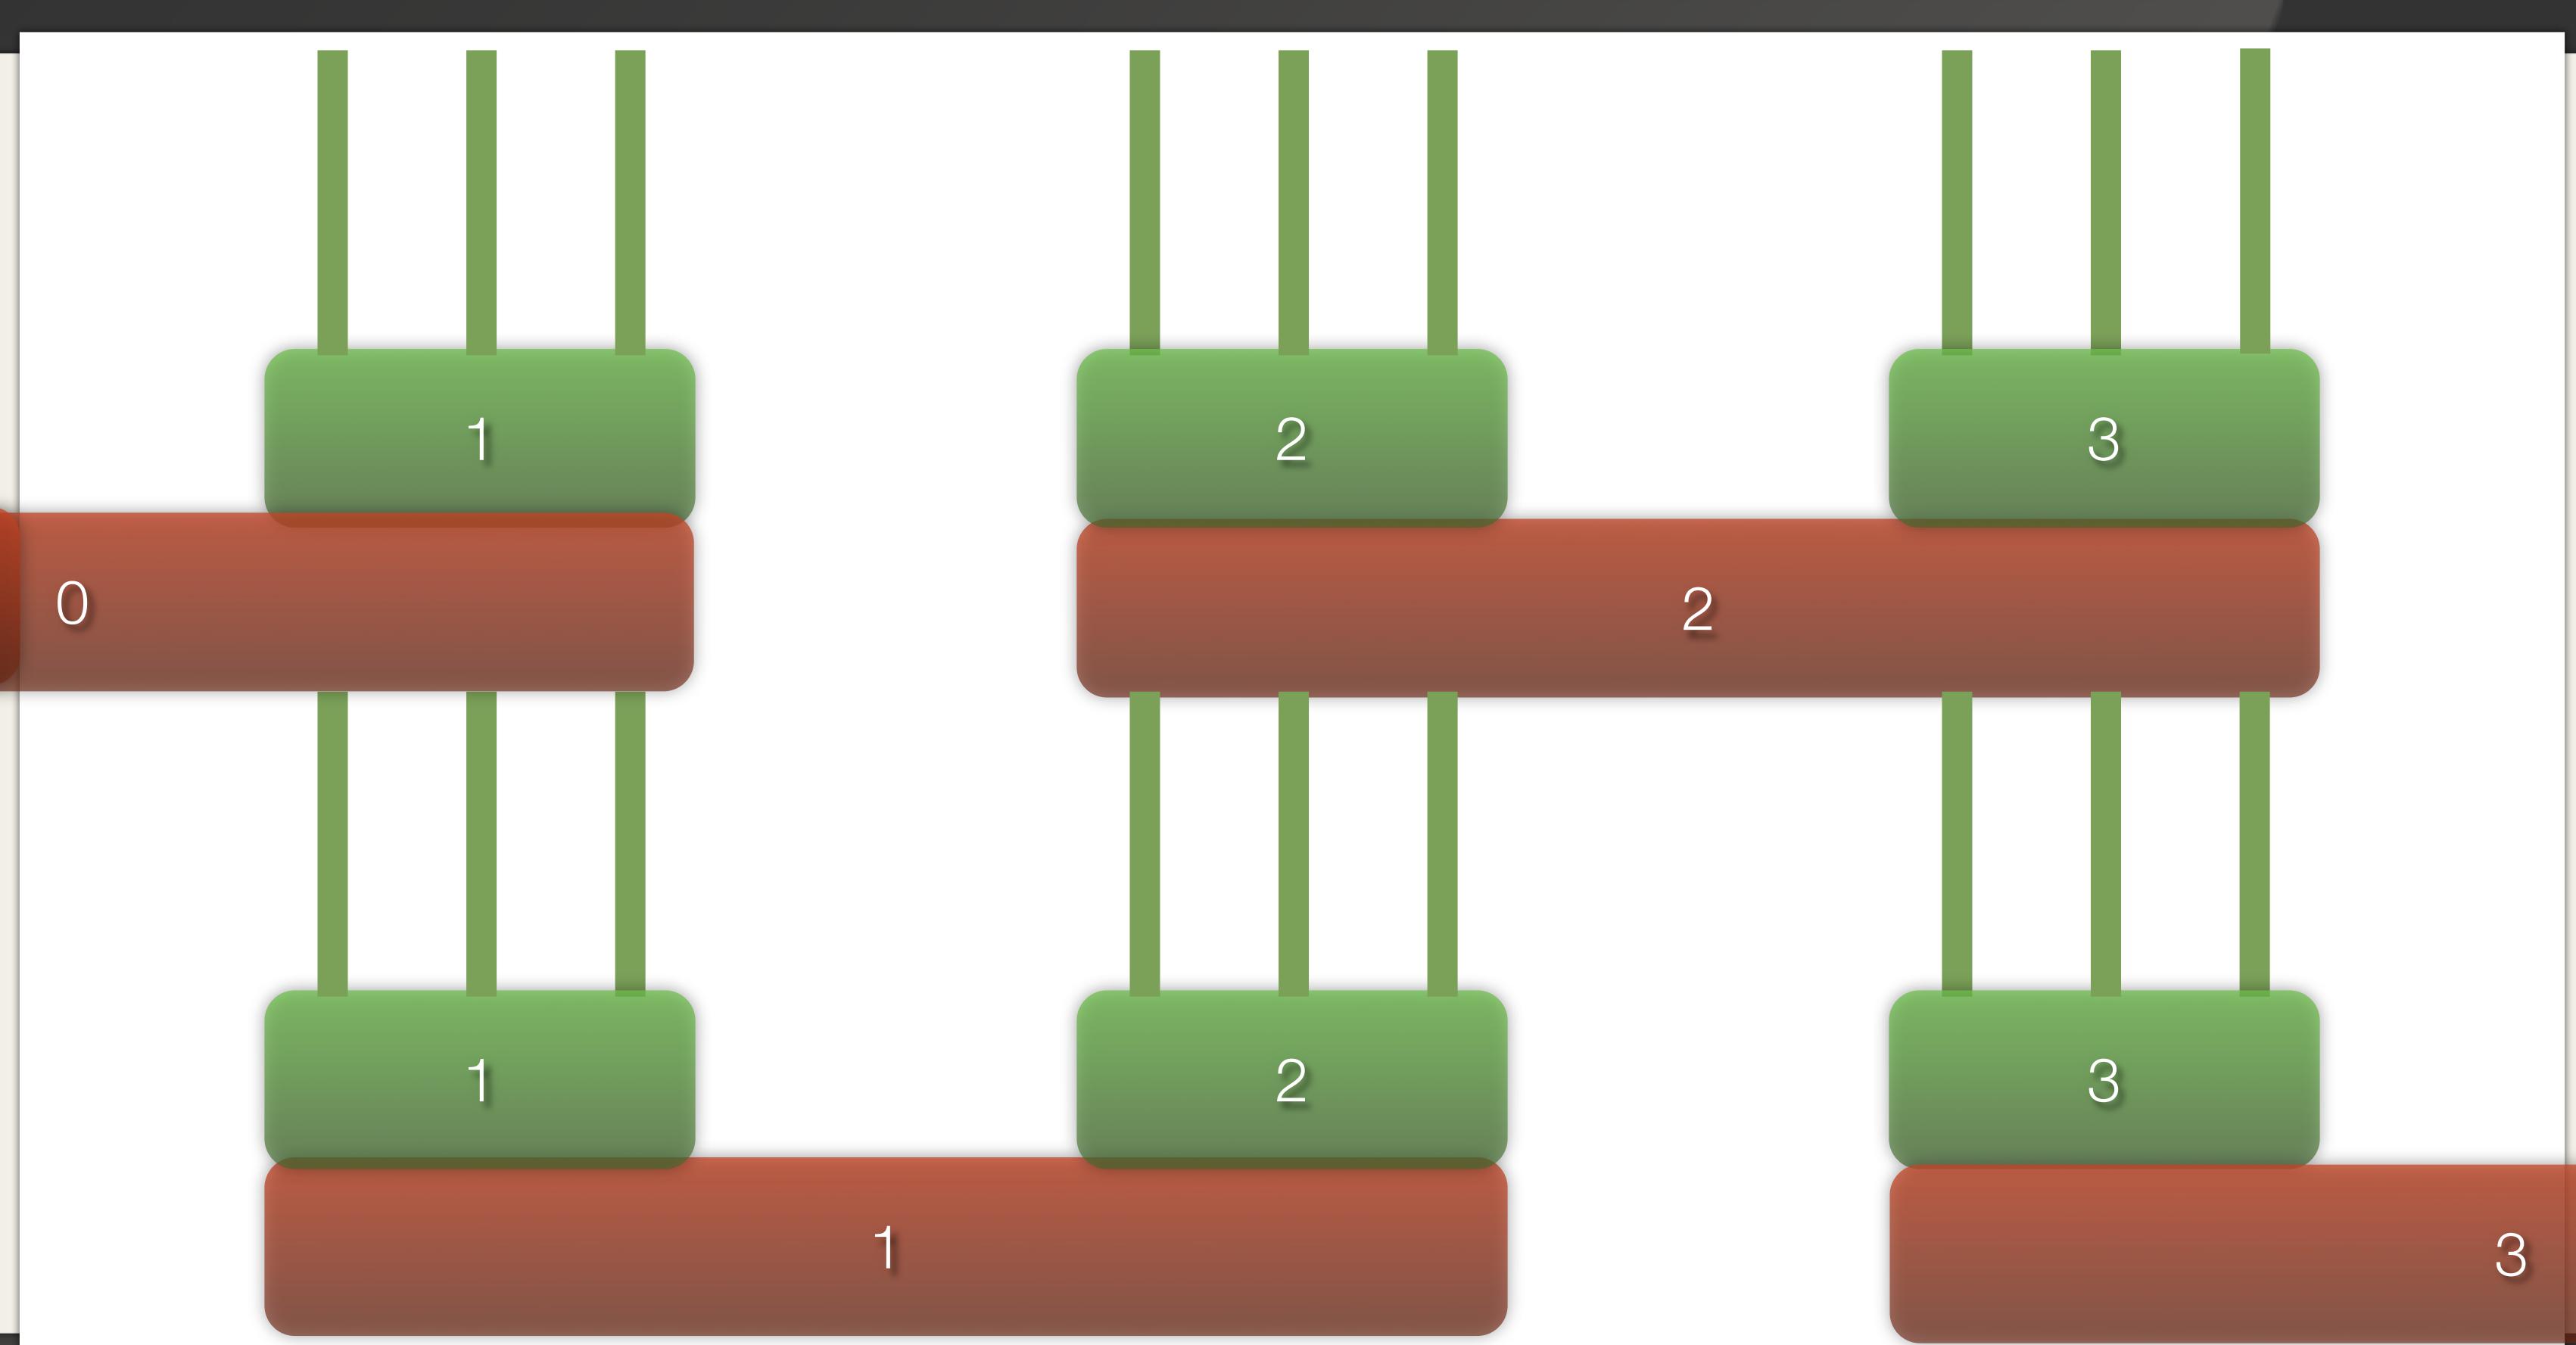
# Application Model

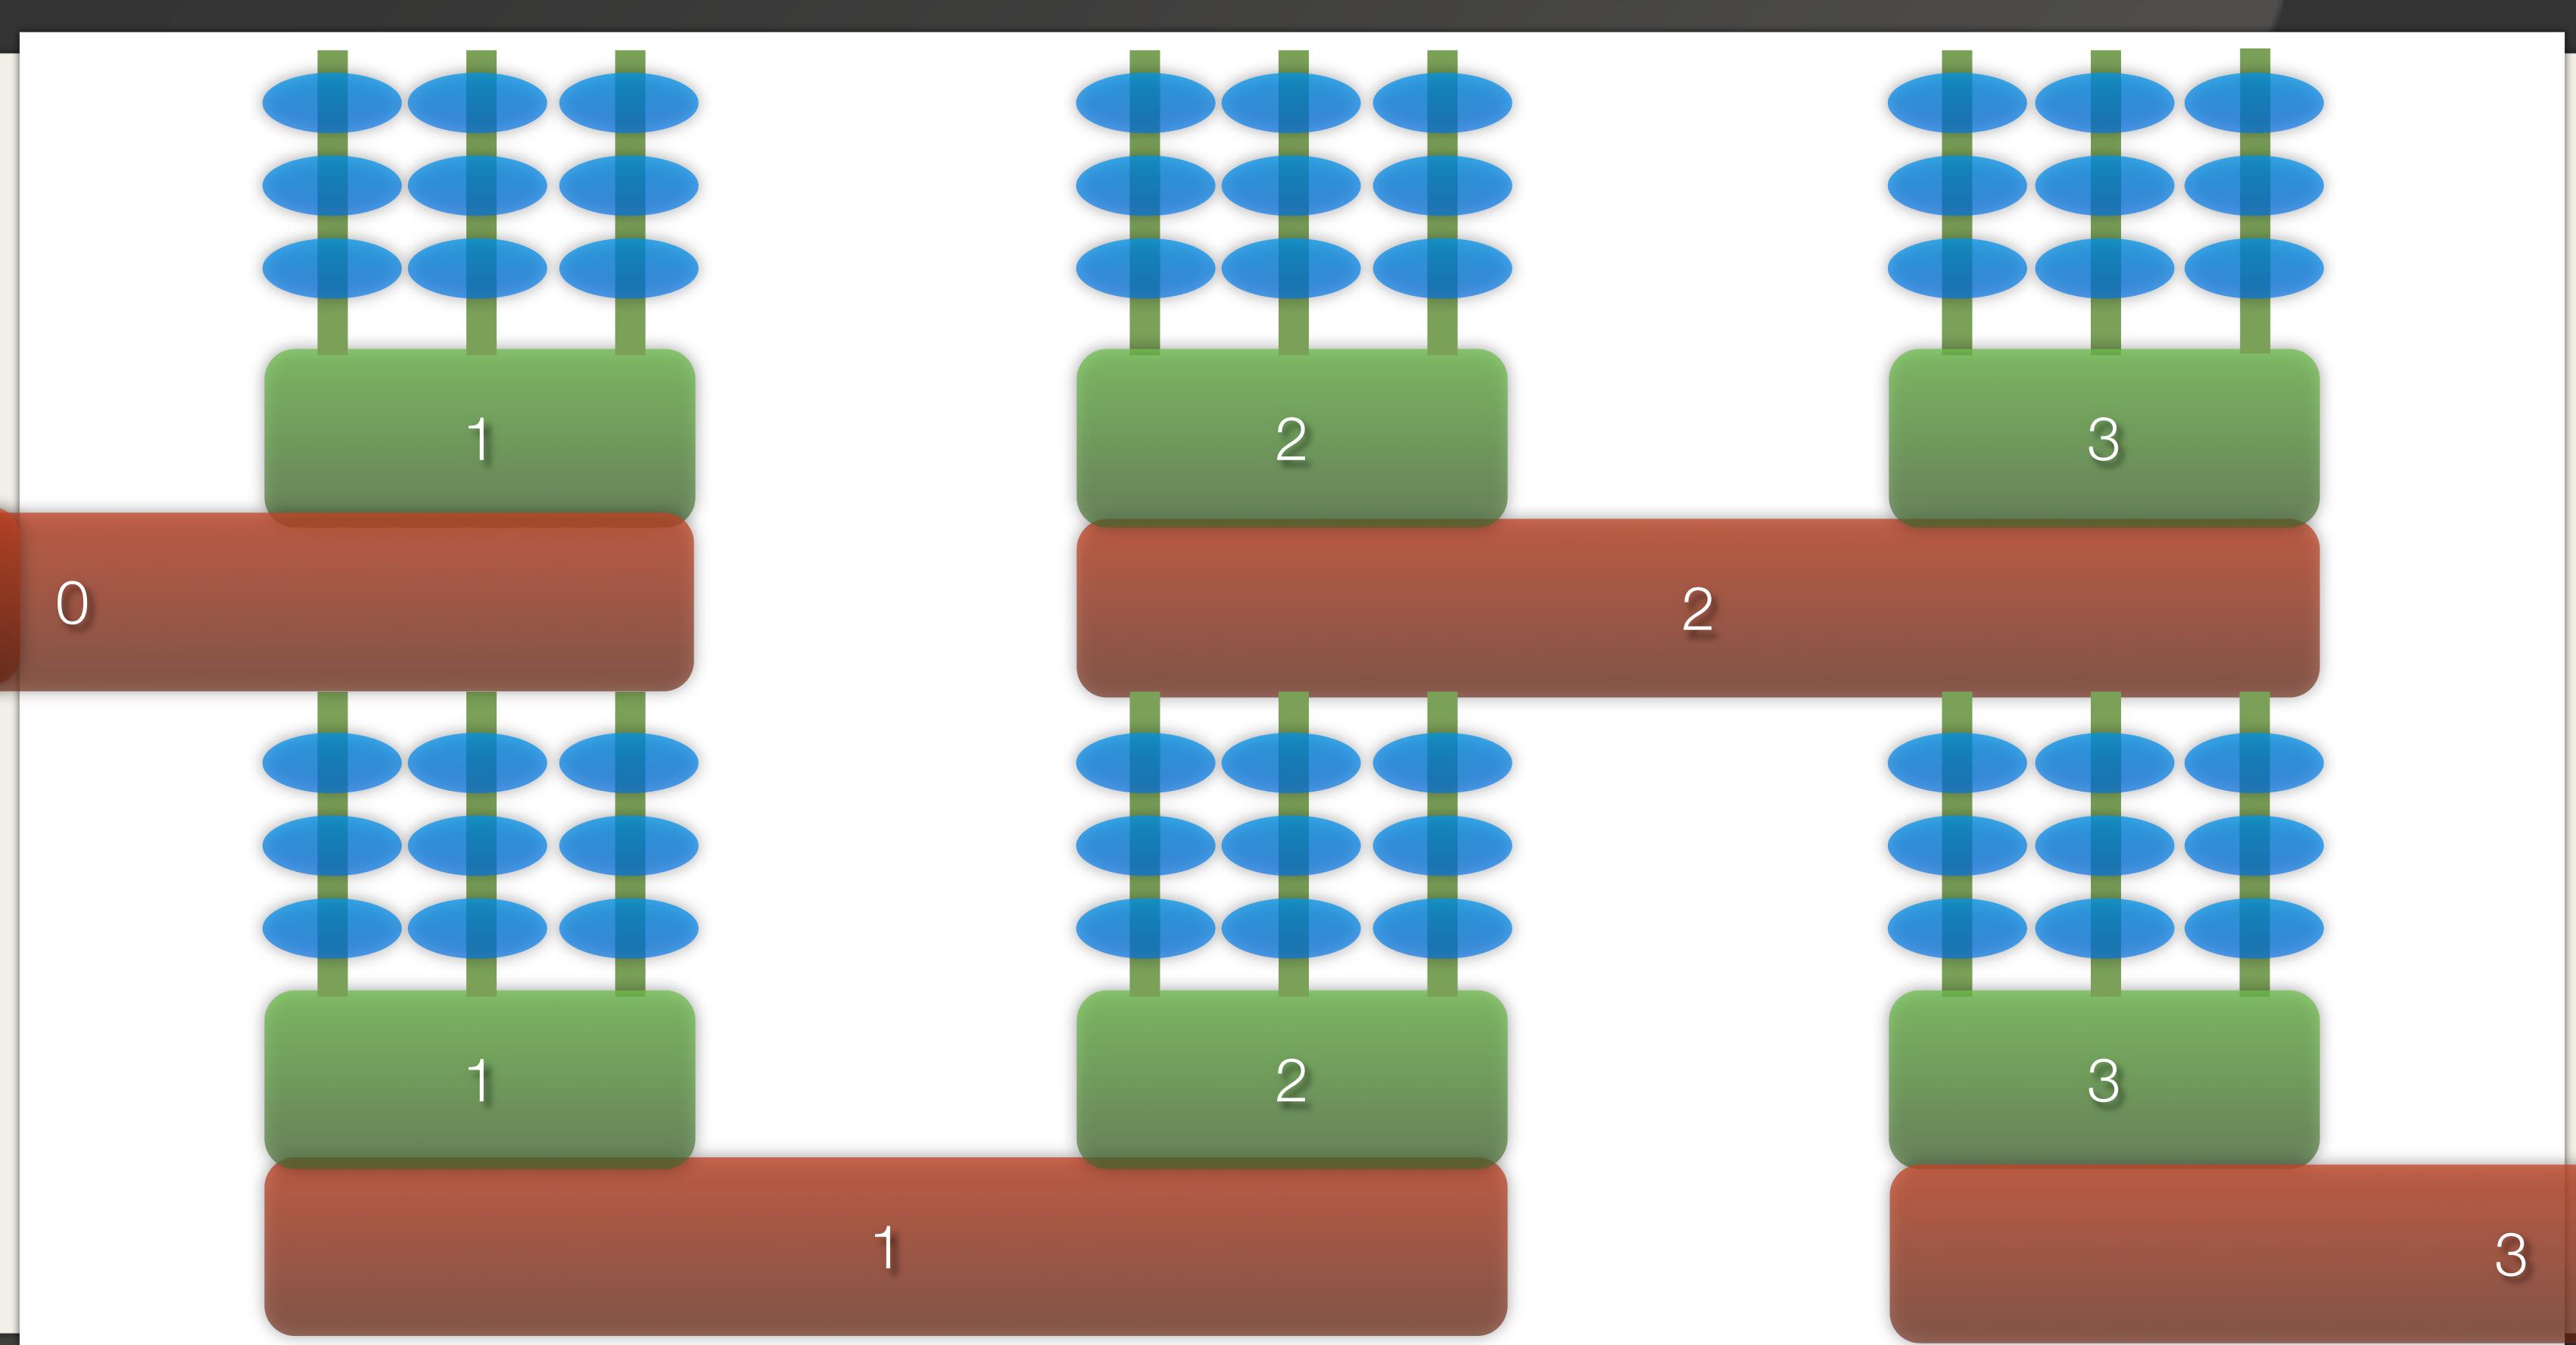
- 'Red' Communicator
  - Includes all processes from own and neighboring energy window
  - Communication:
    - 1x MPI\_Scatter
    - 2x (MPI\_Send\_recv\_replace)
    - Many MPI\_Send & MPI\_Recv

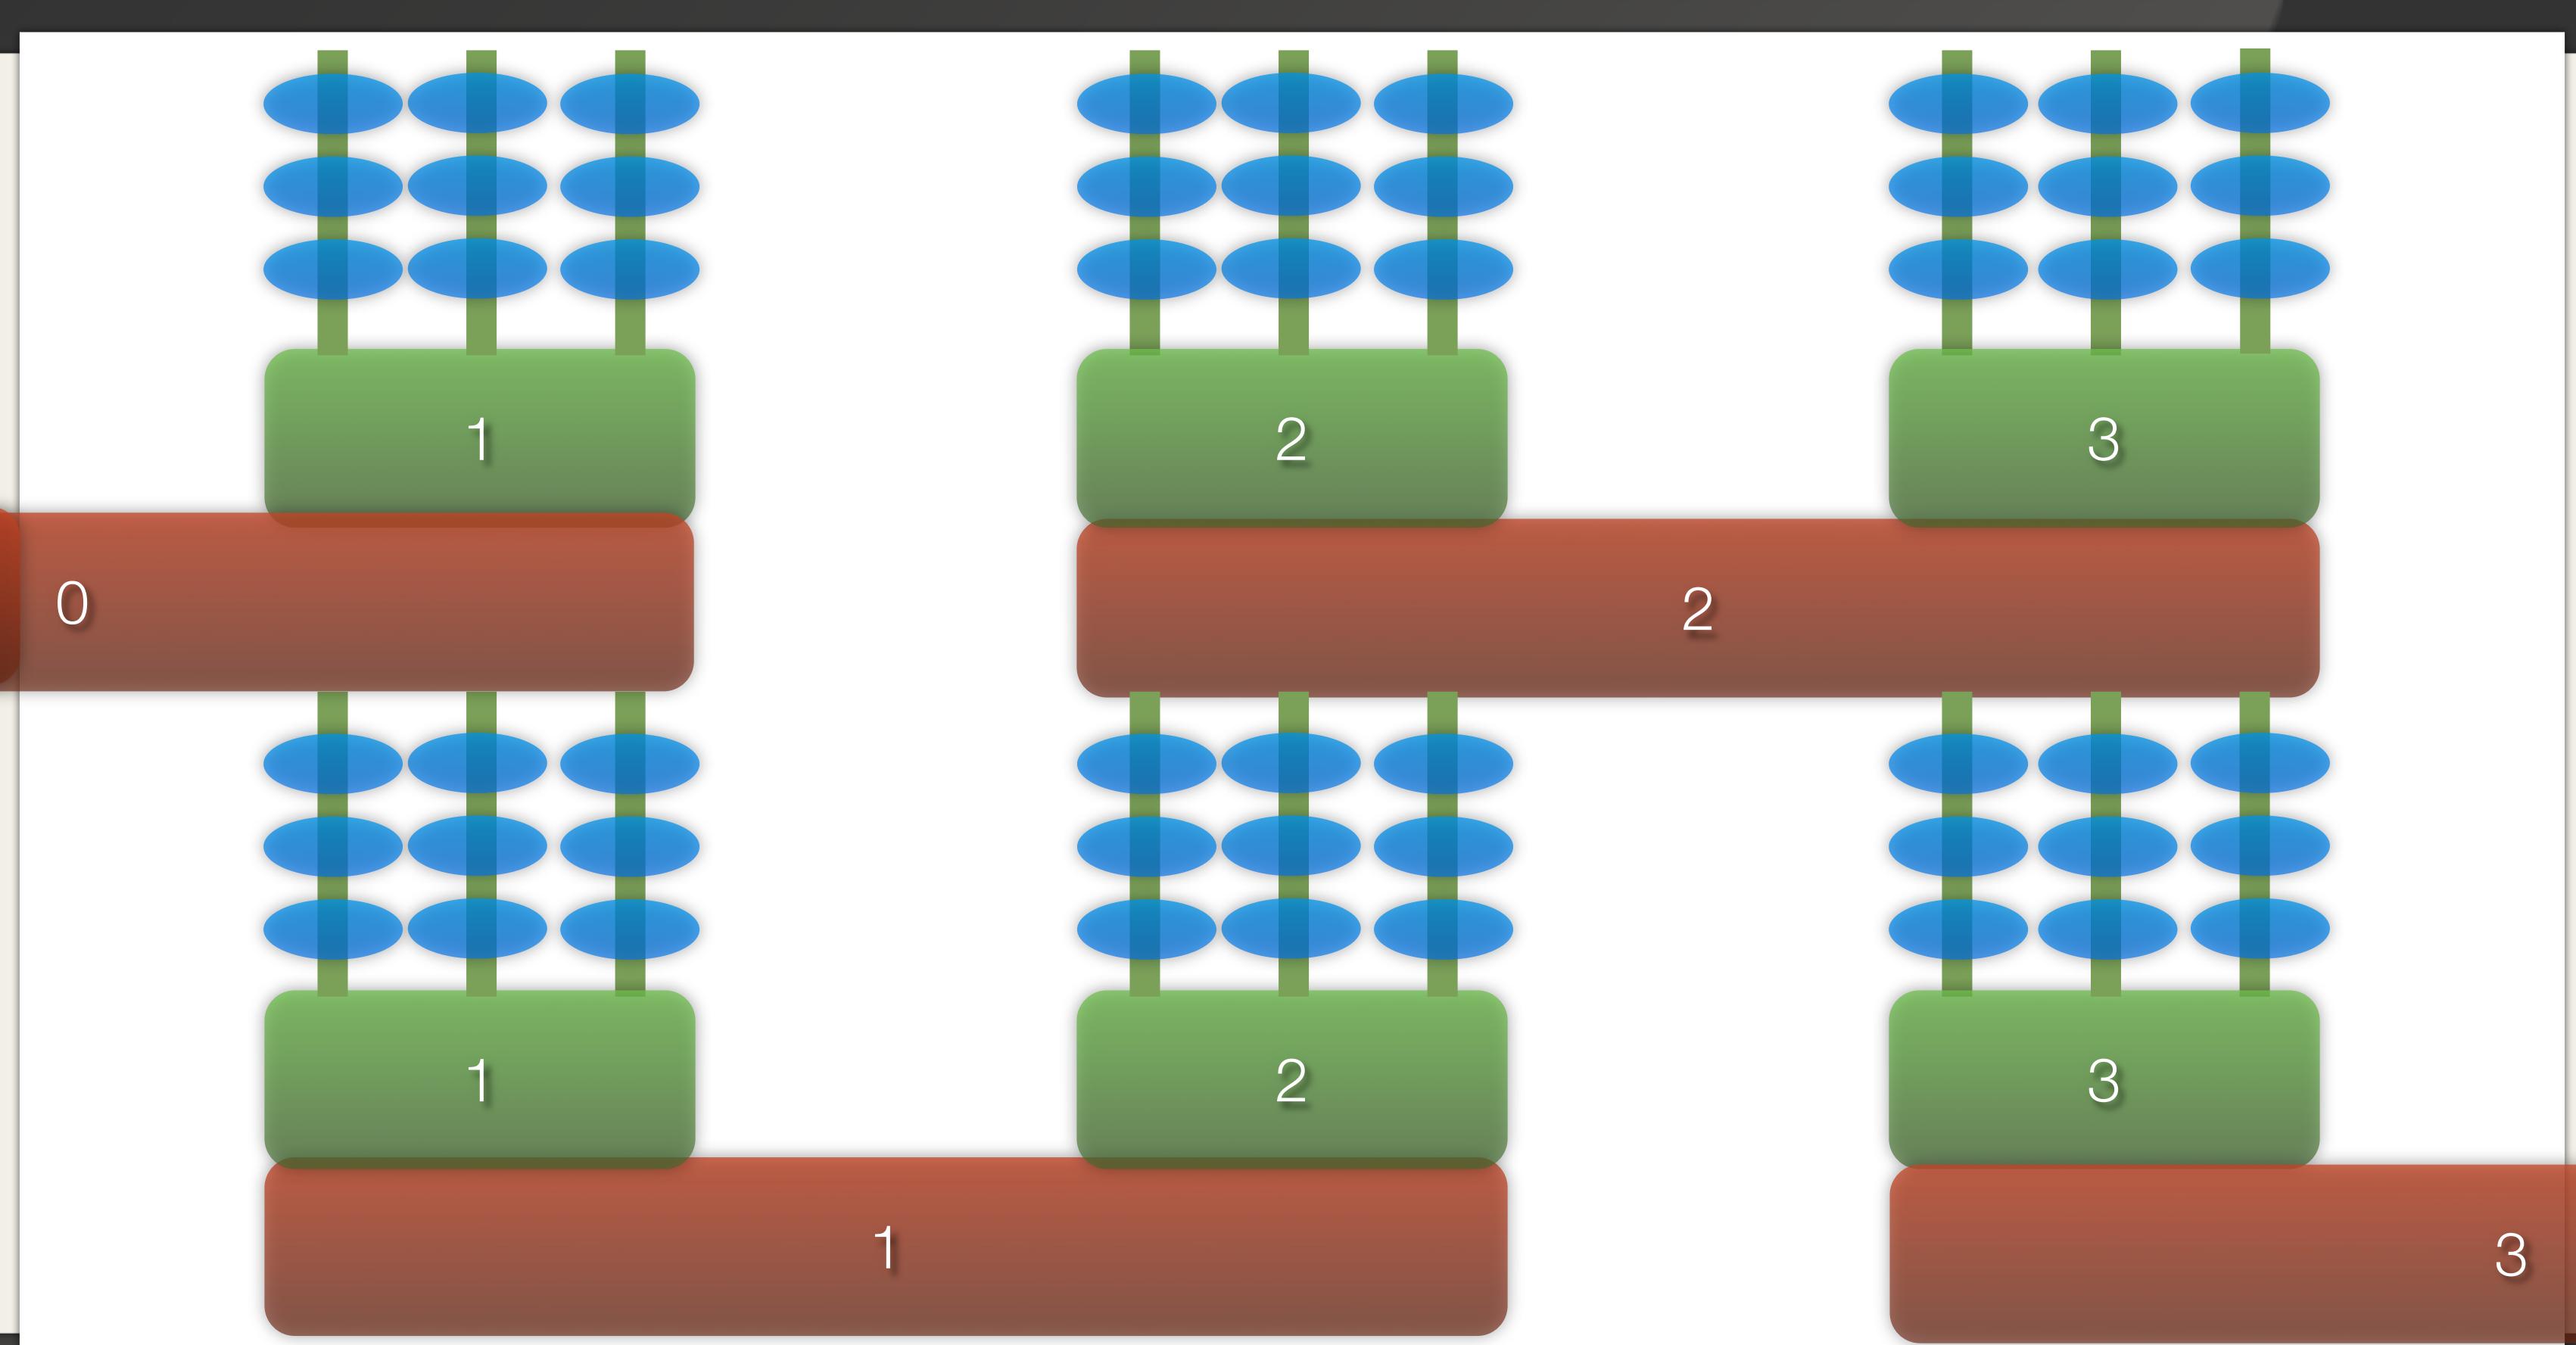


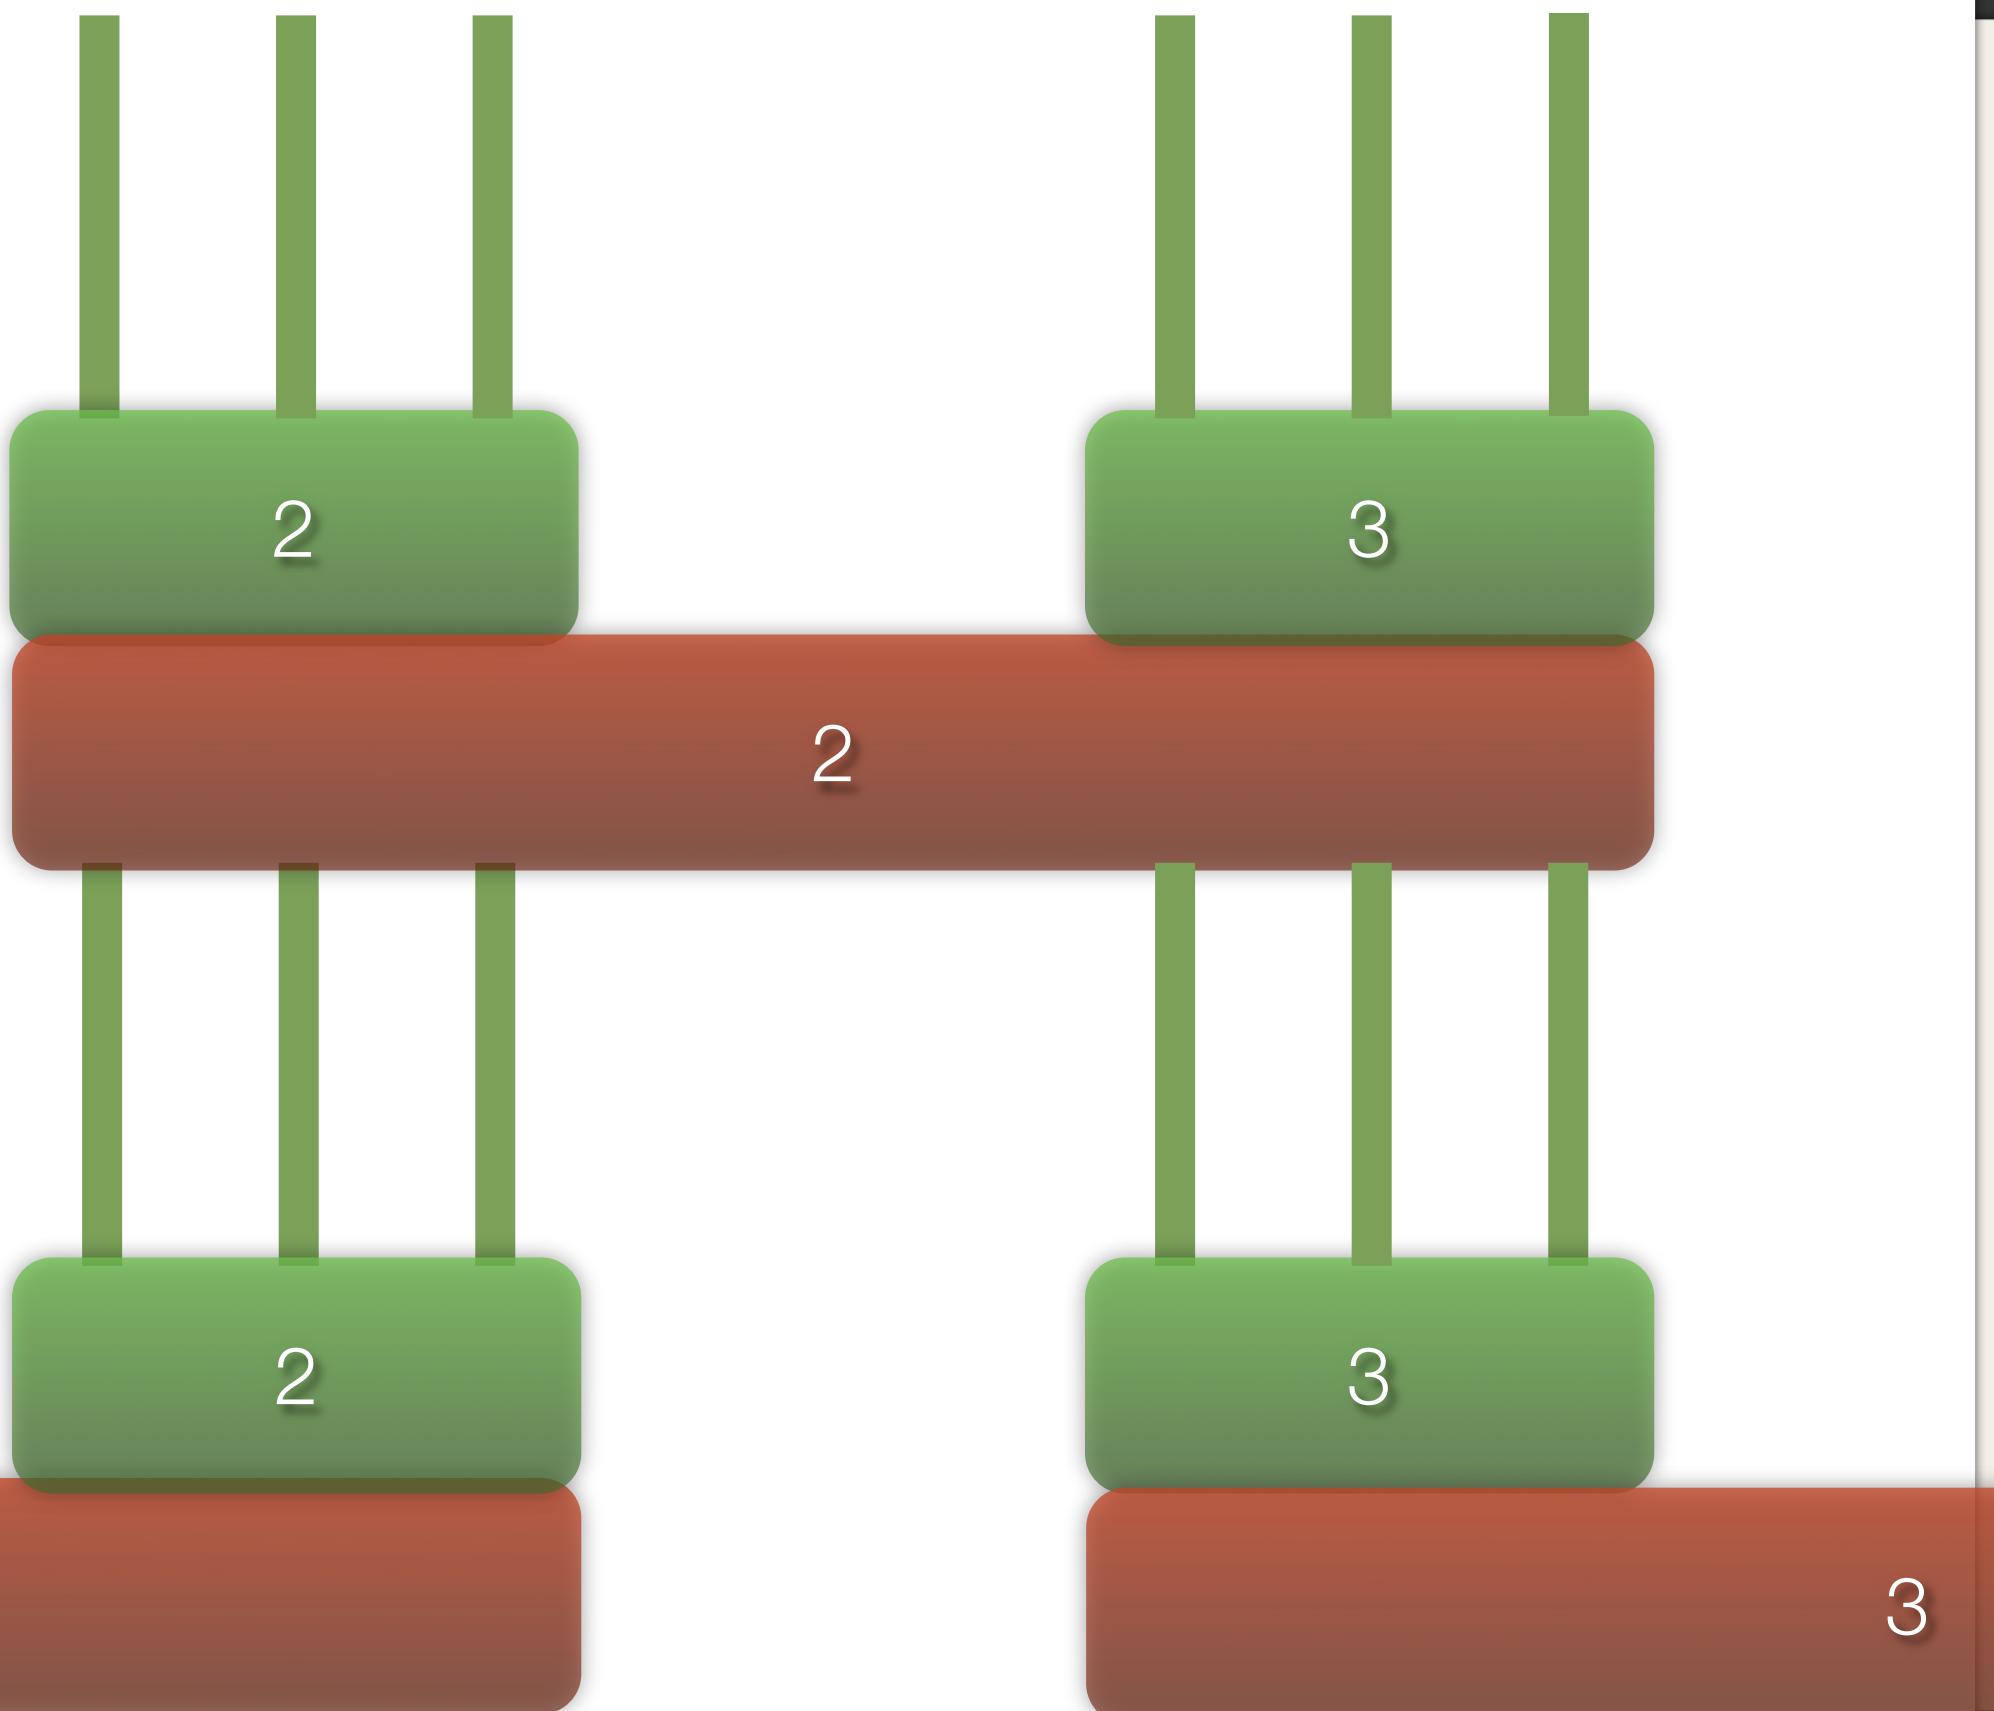
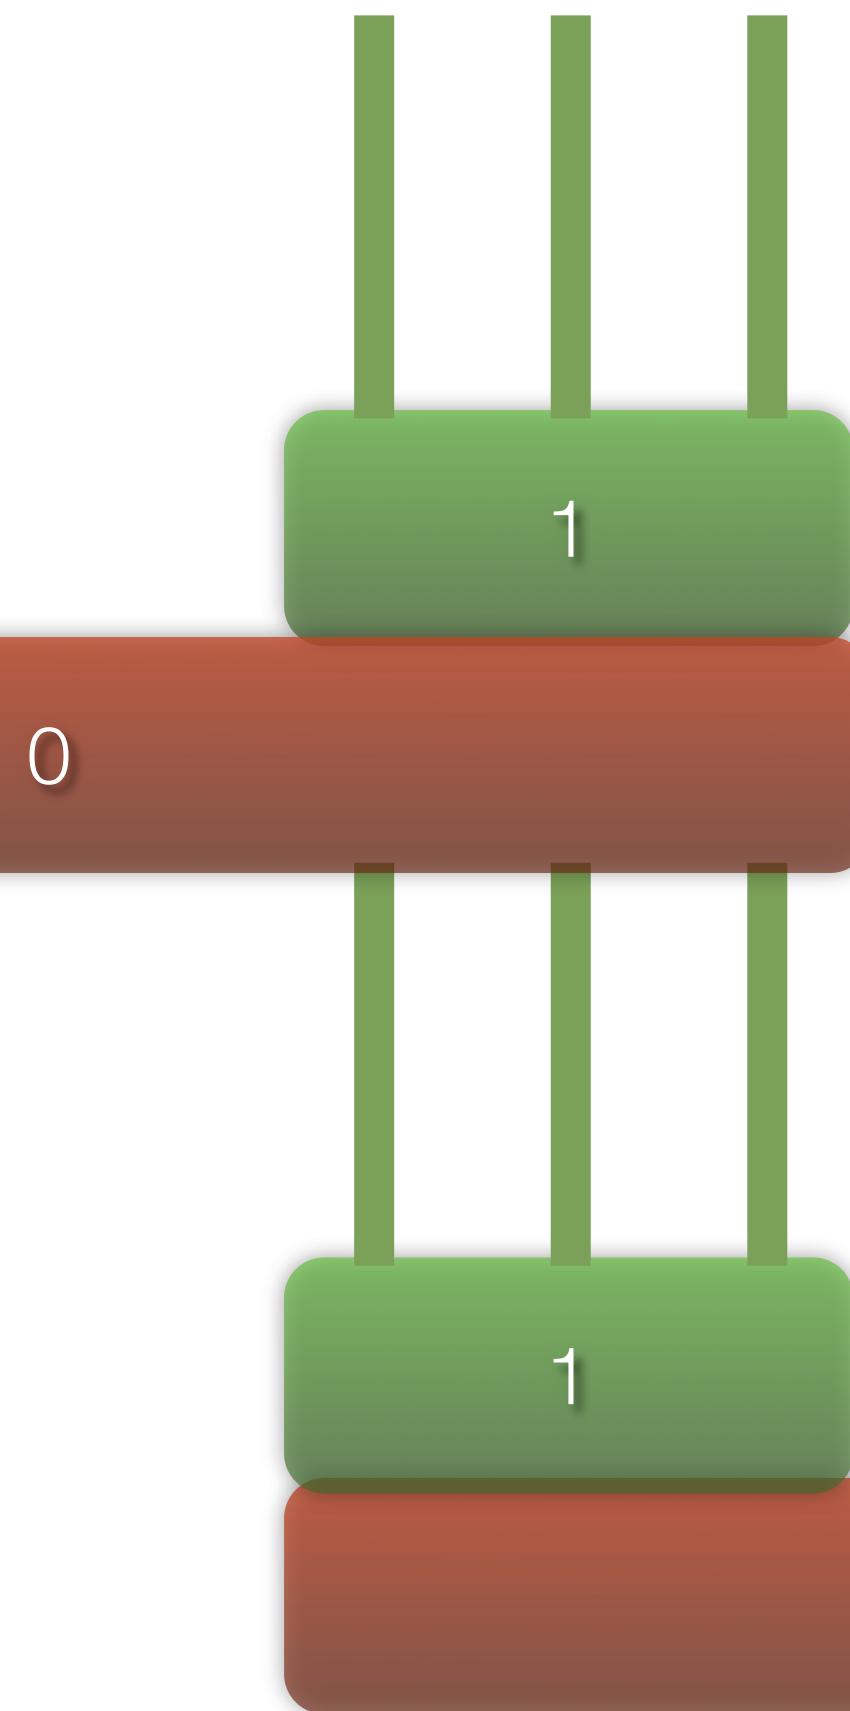
# Recovery Model

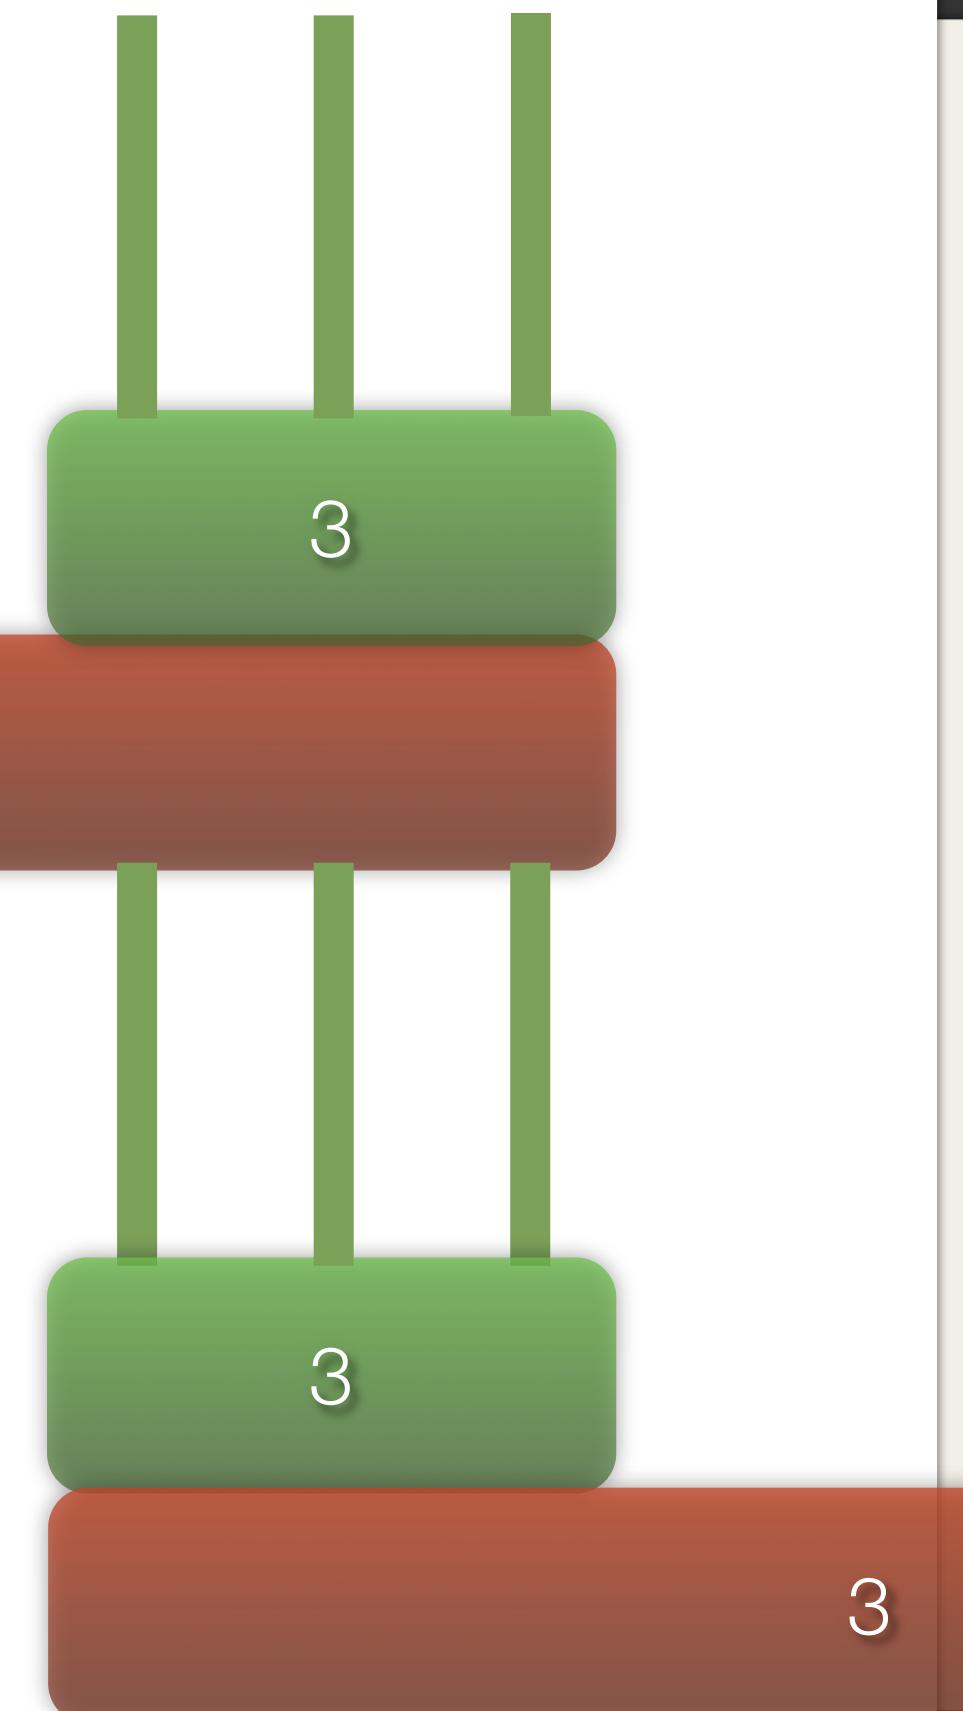
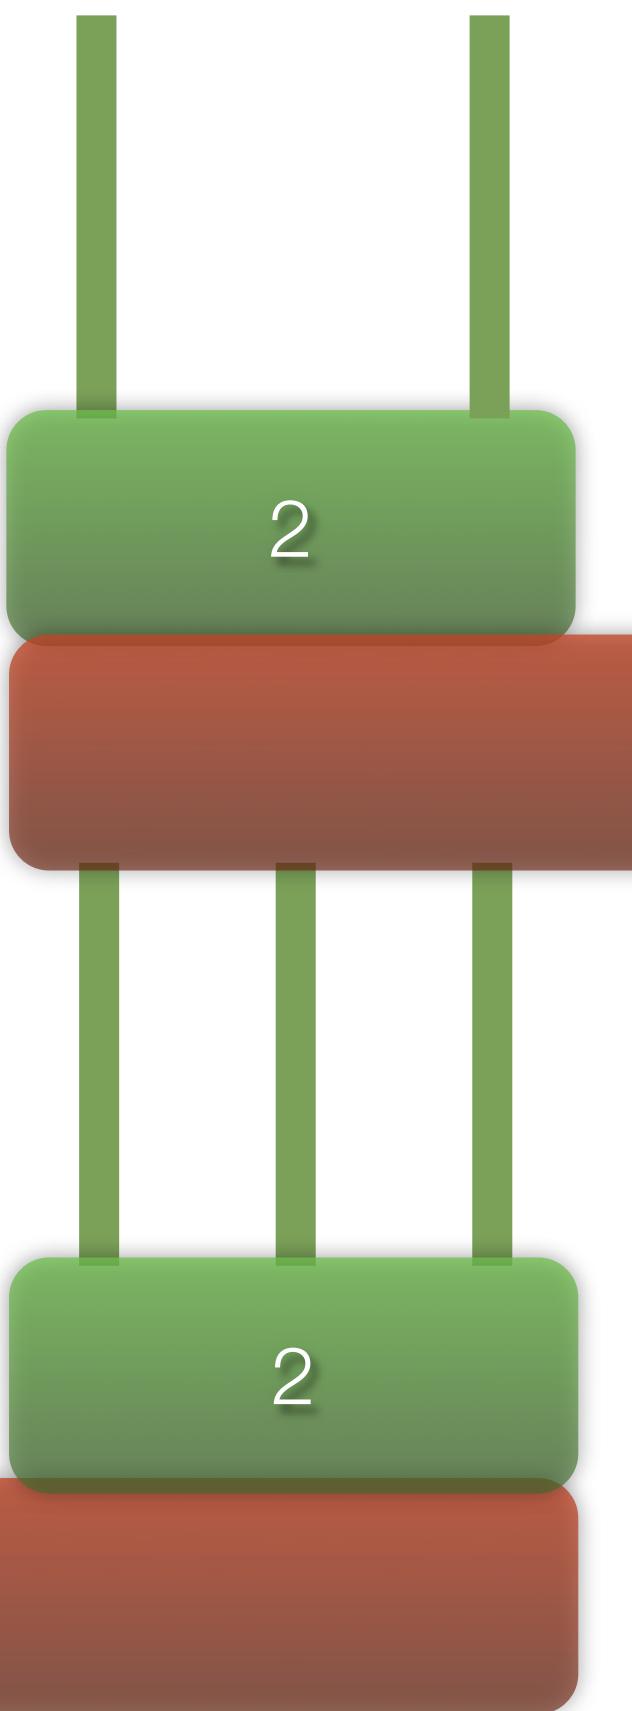
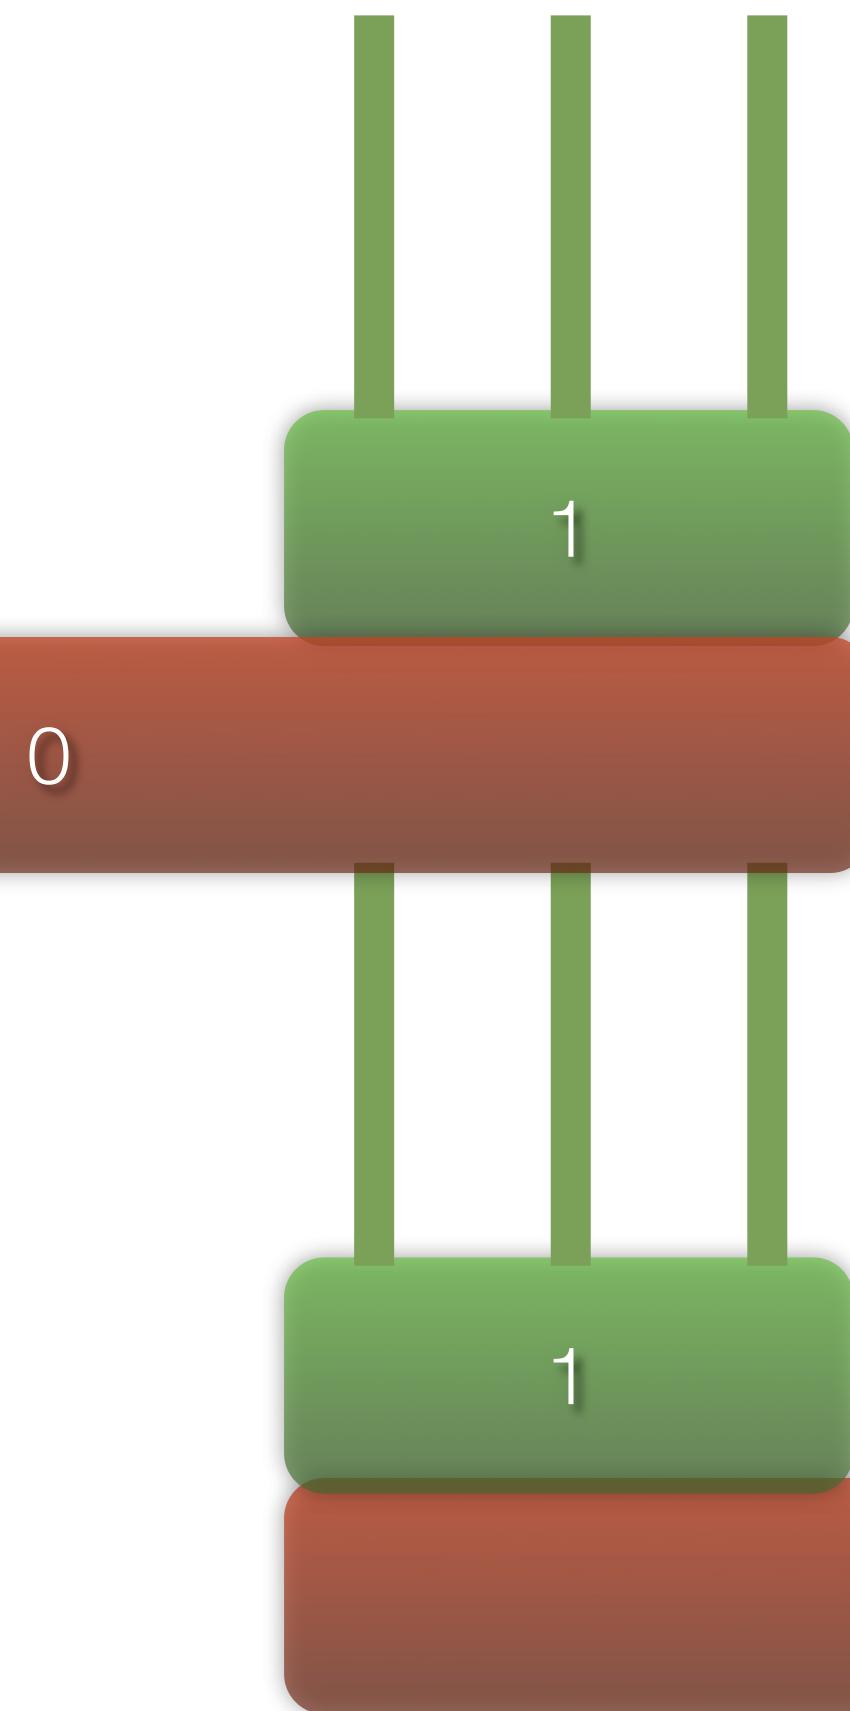


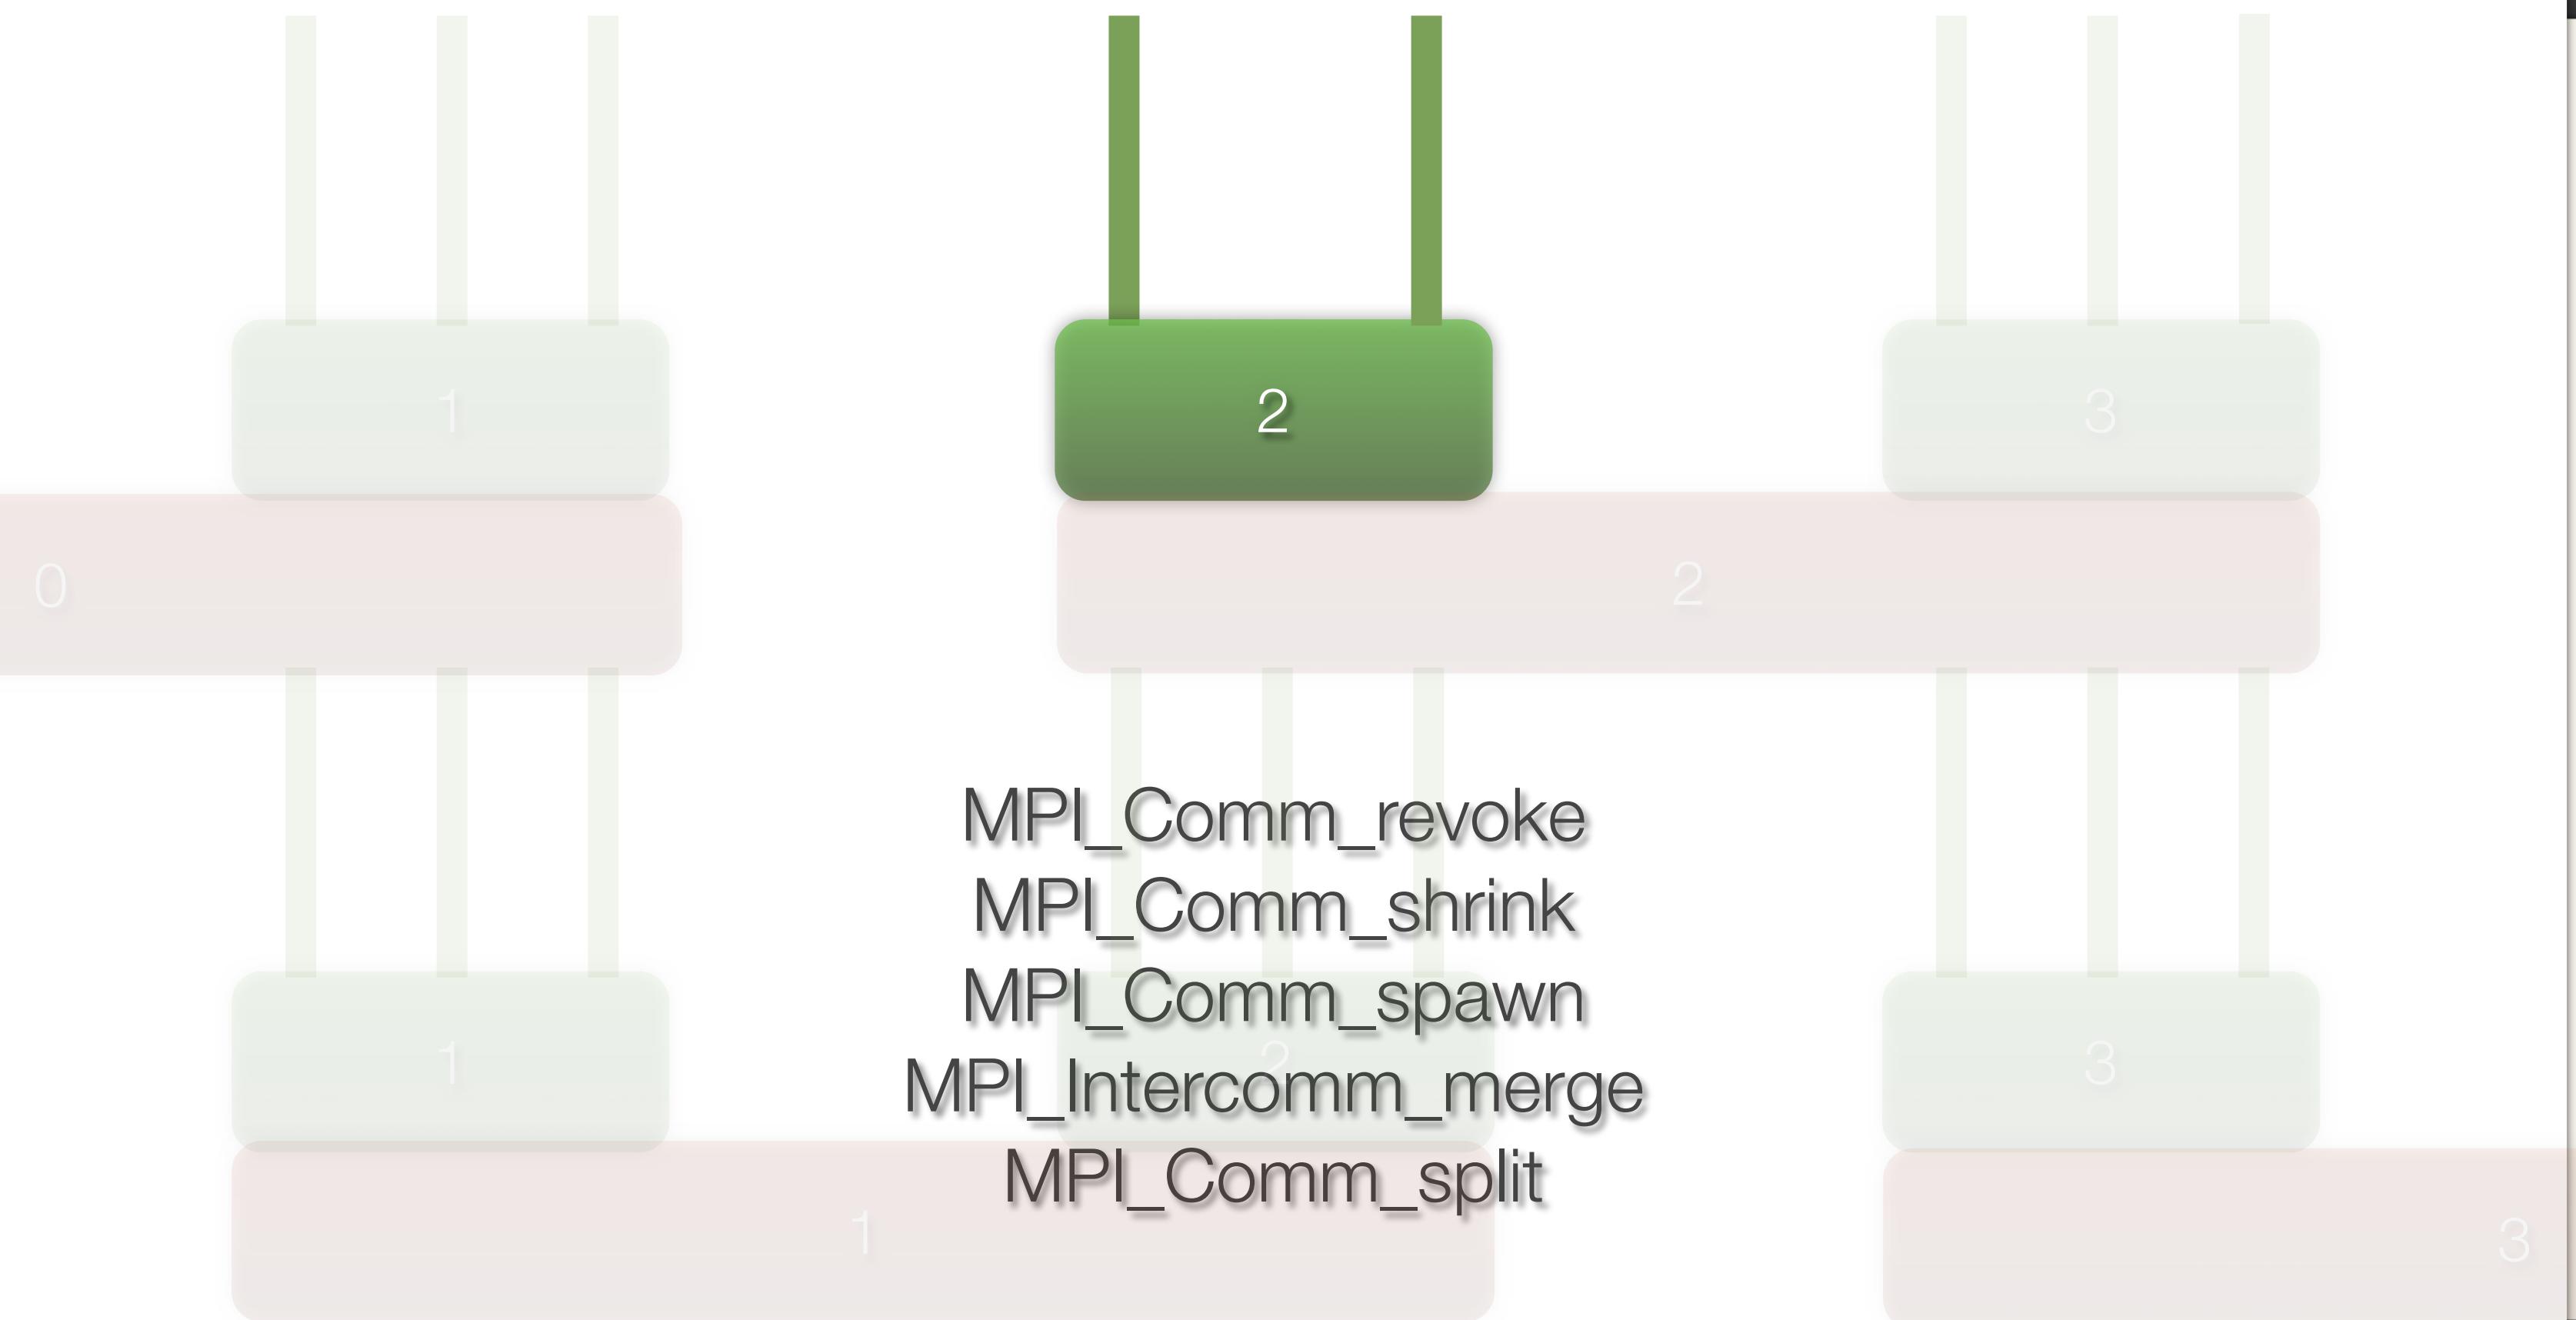


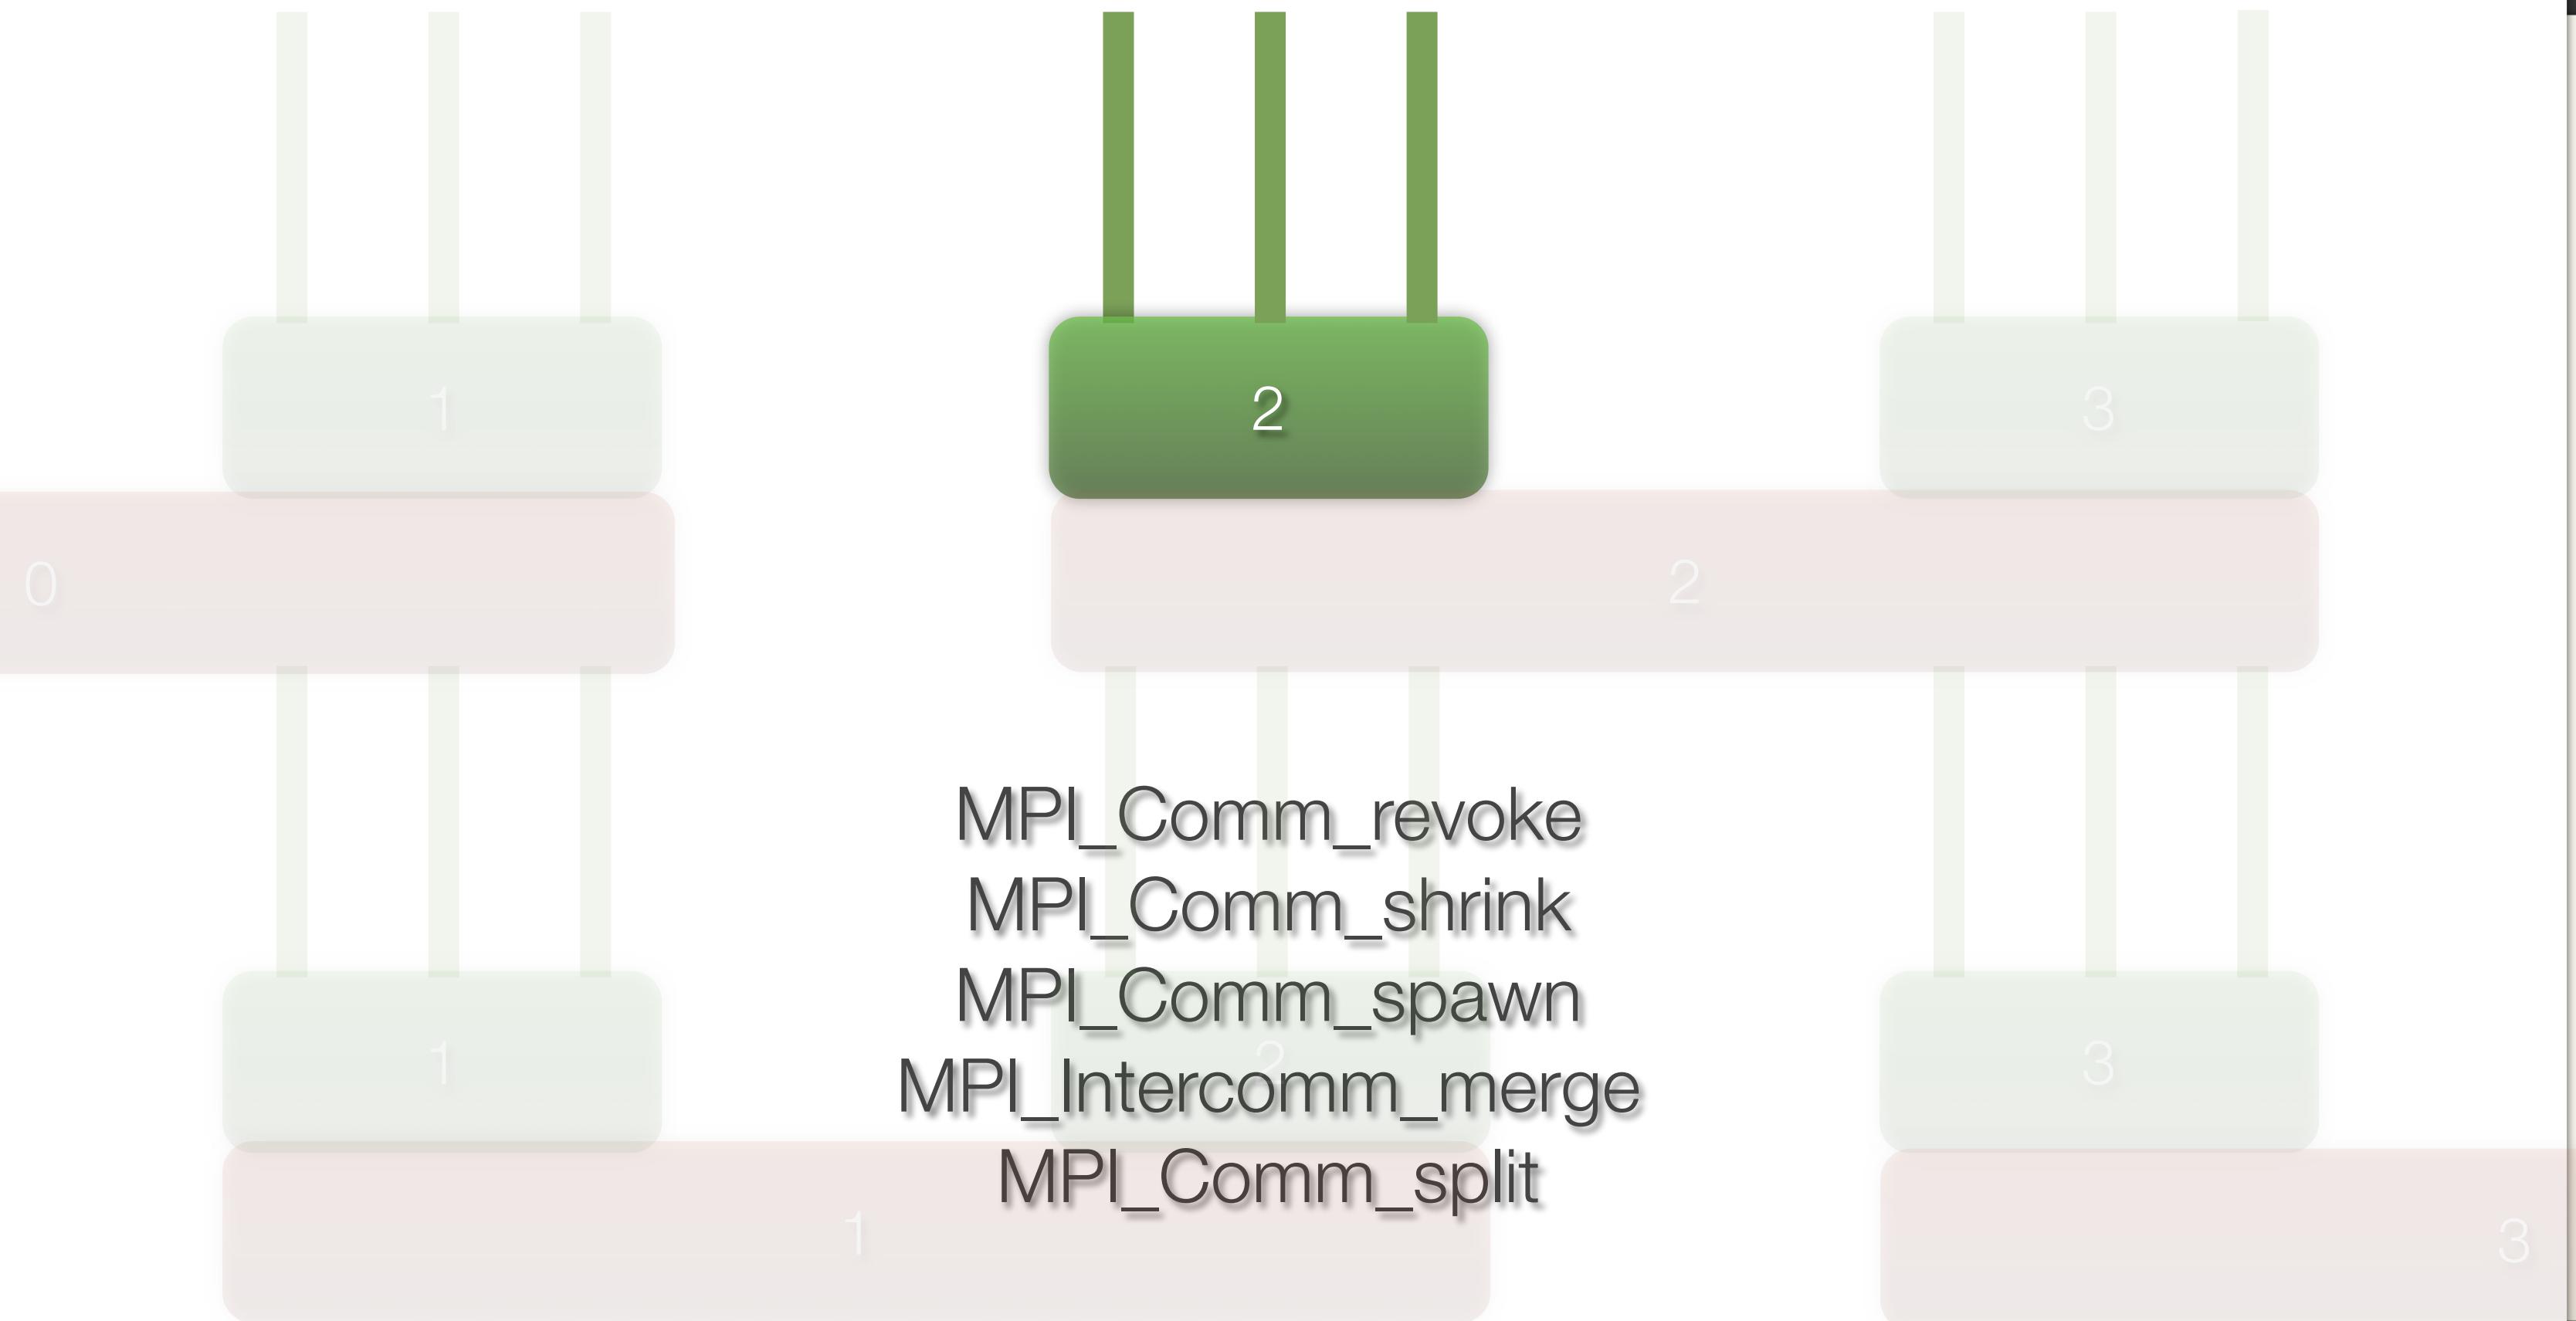


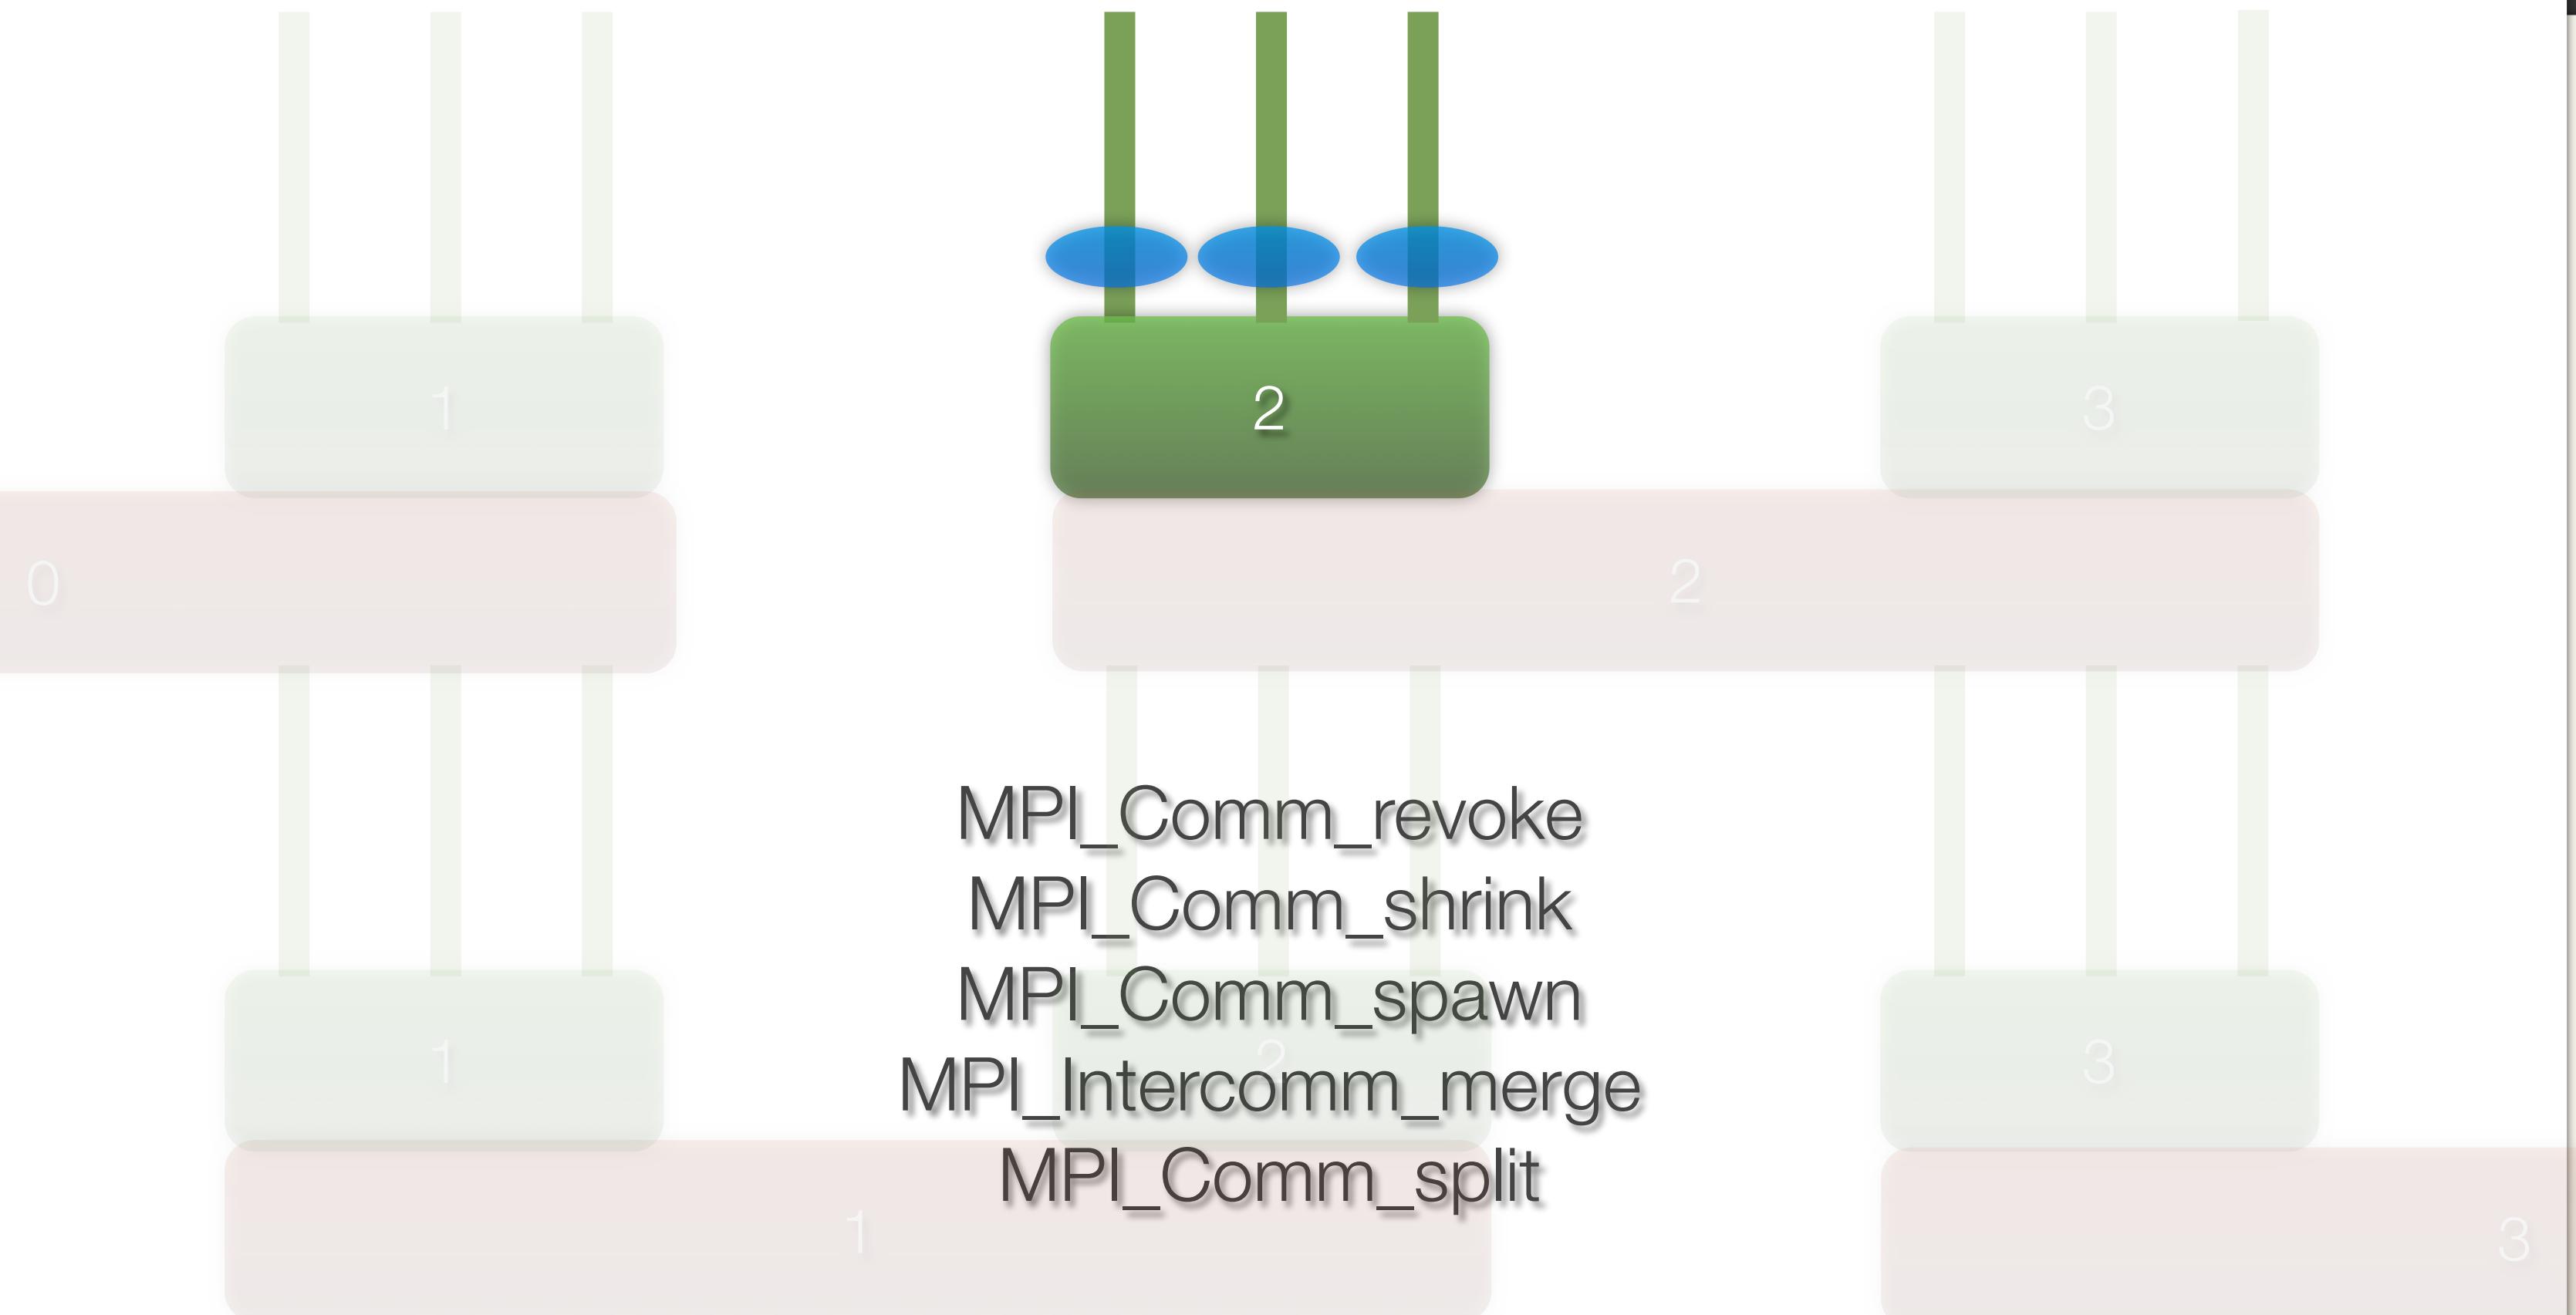


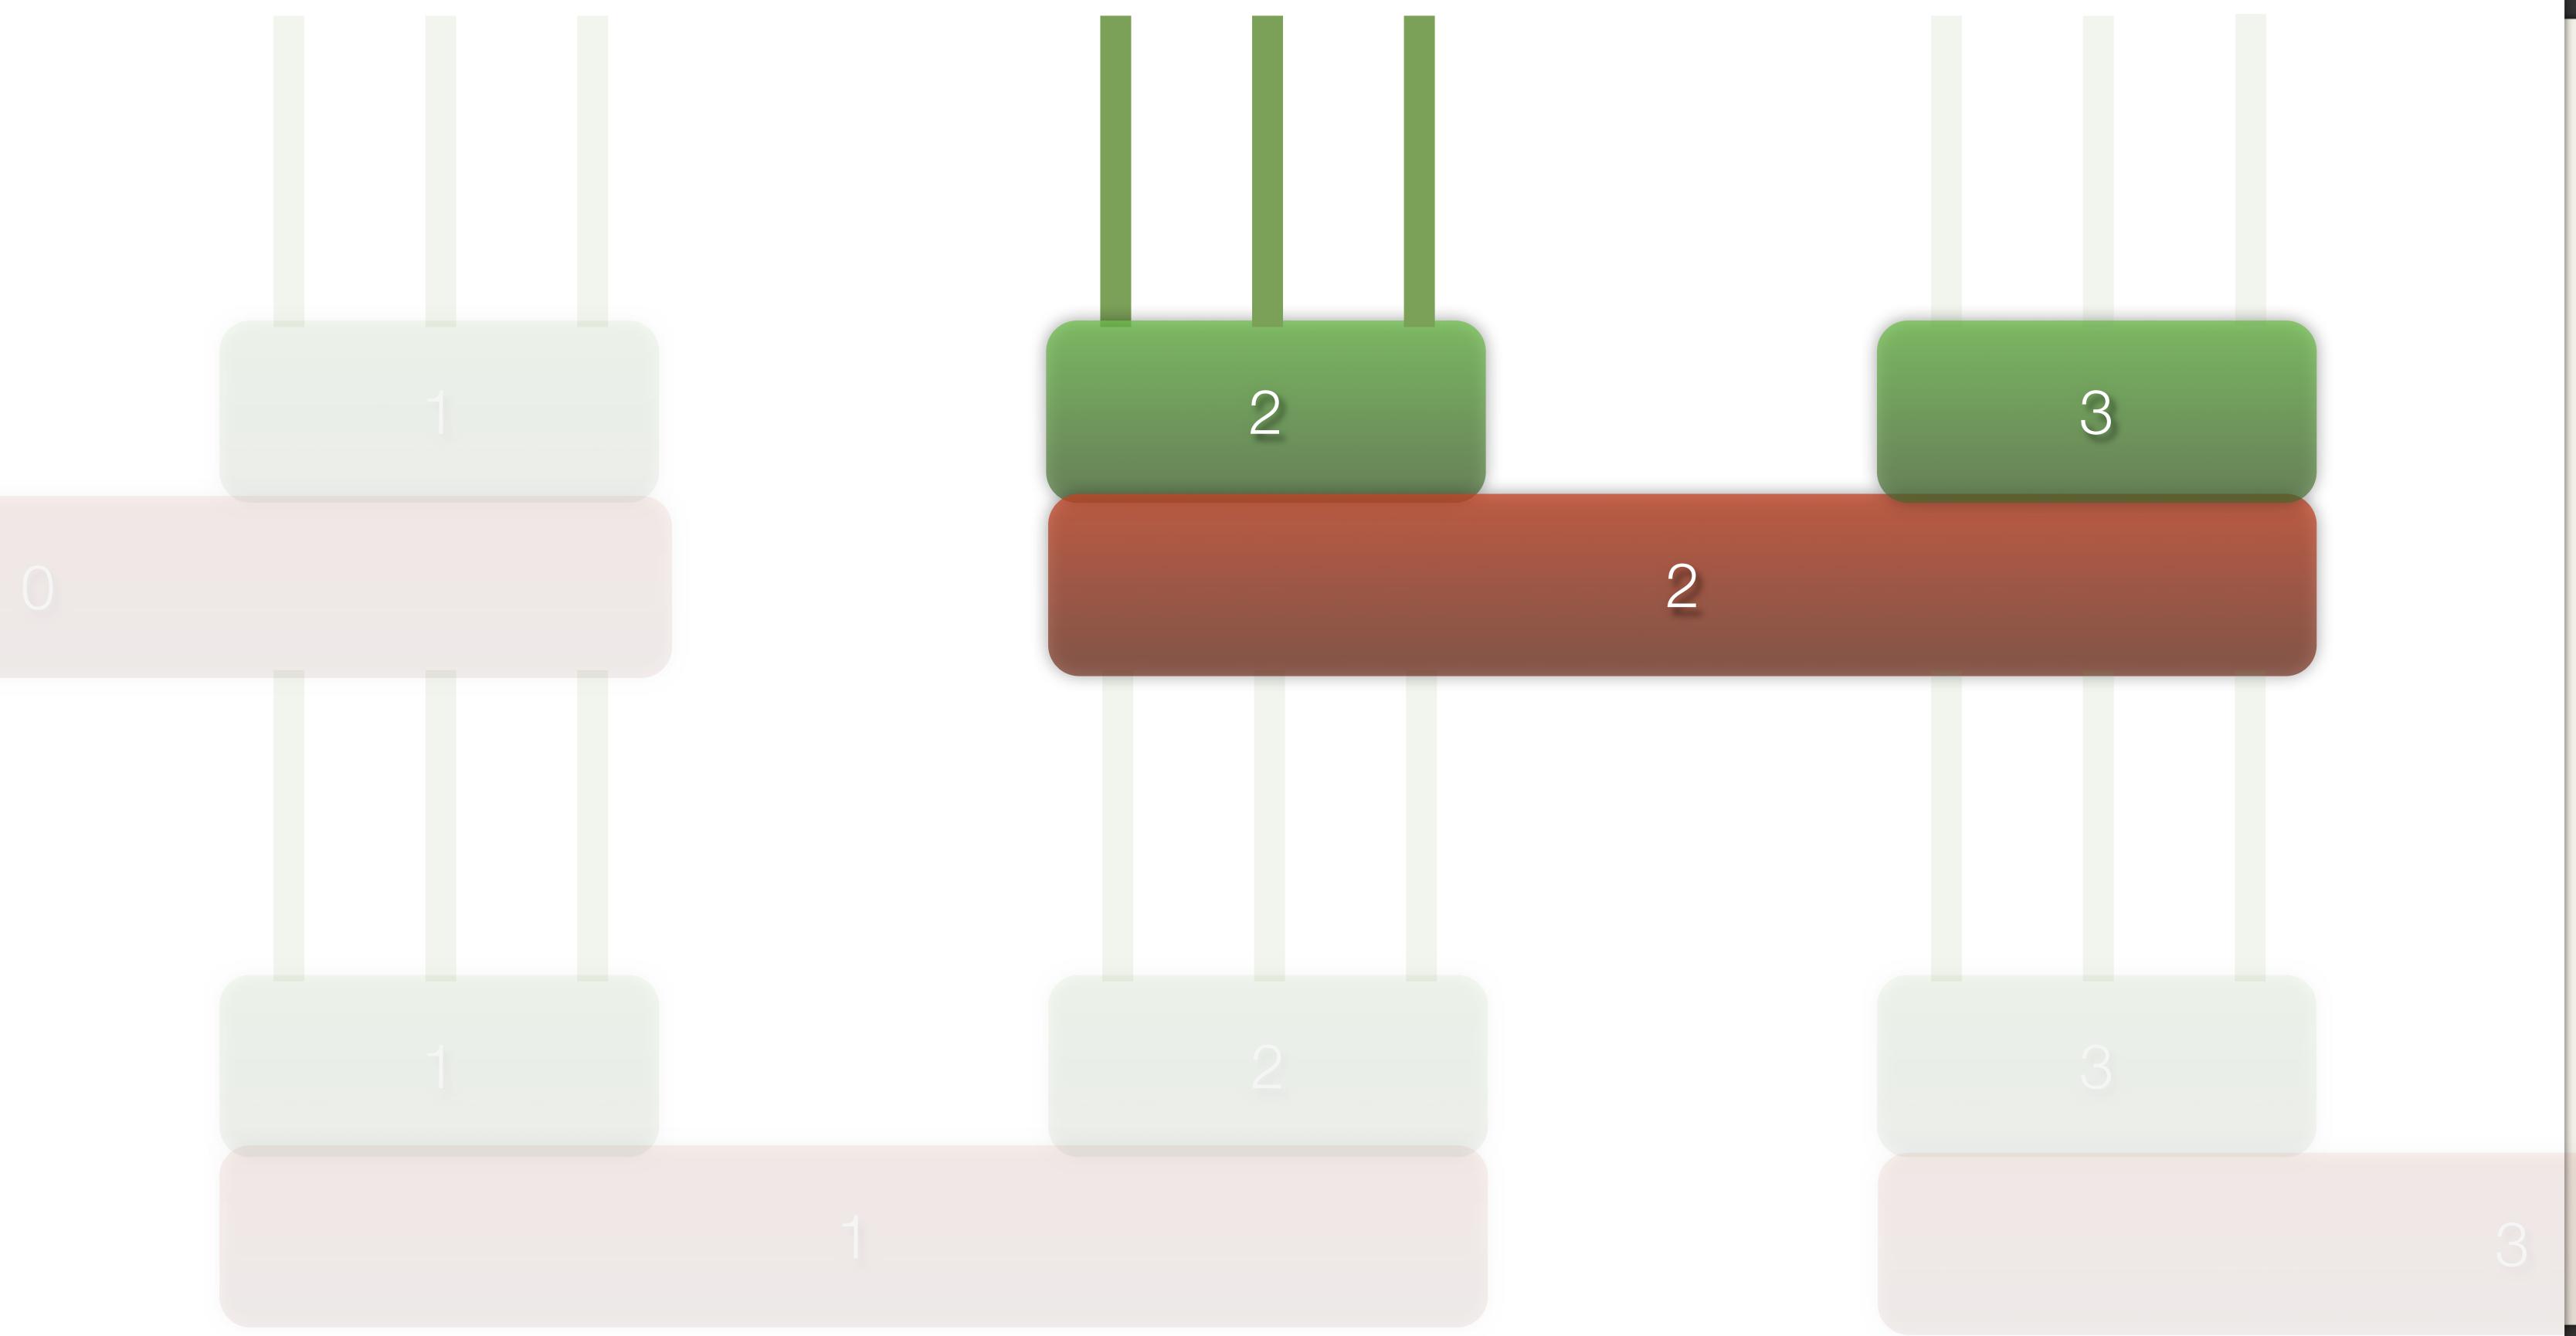


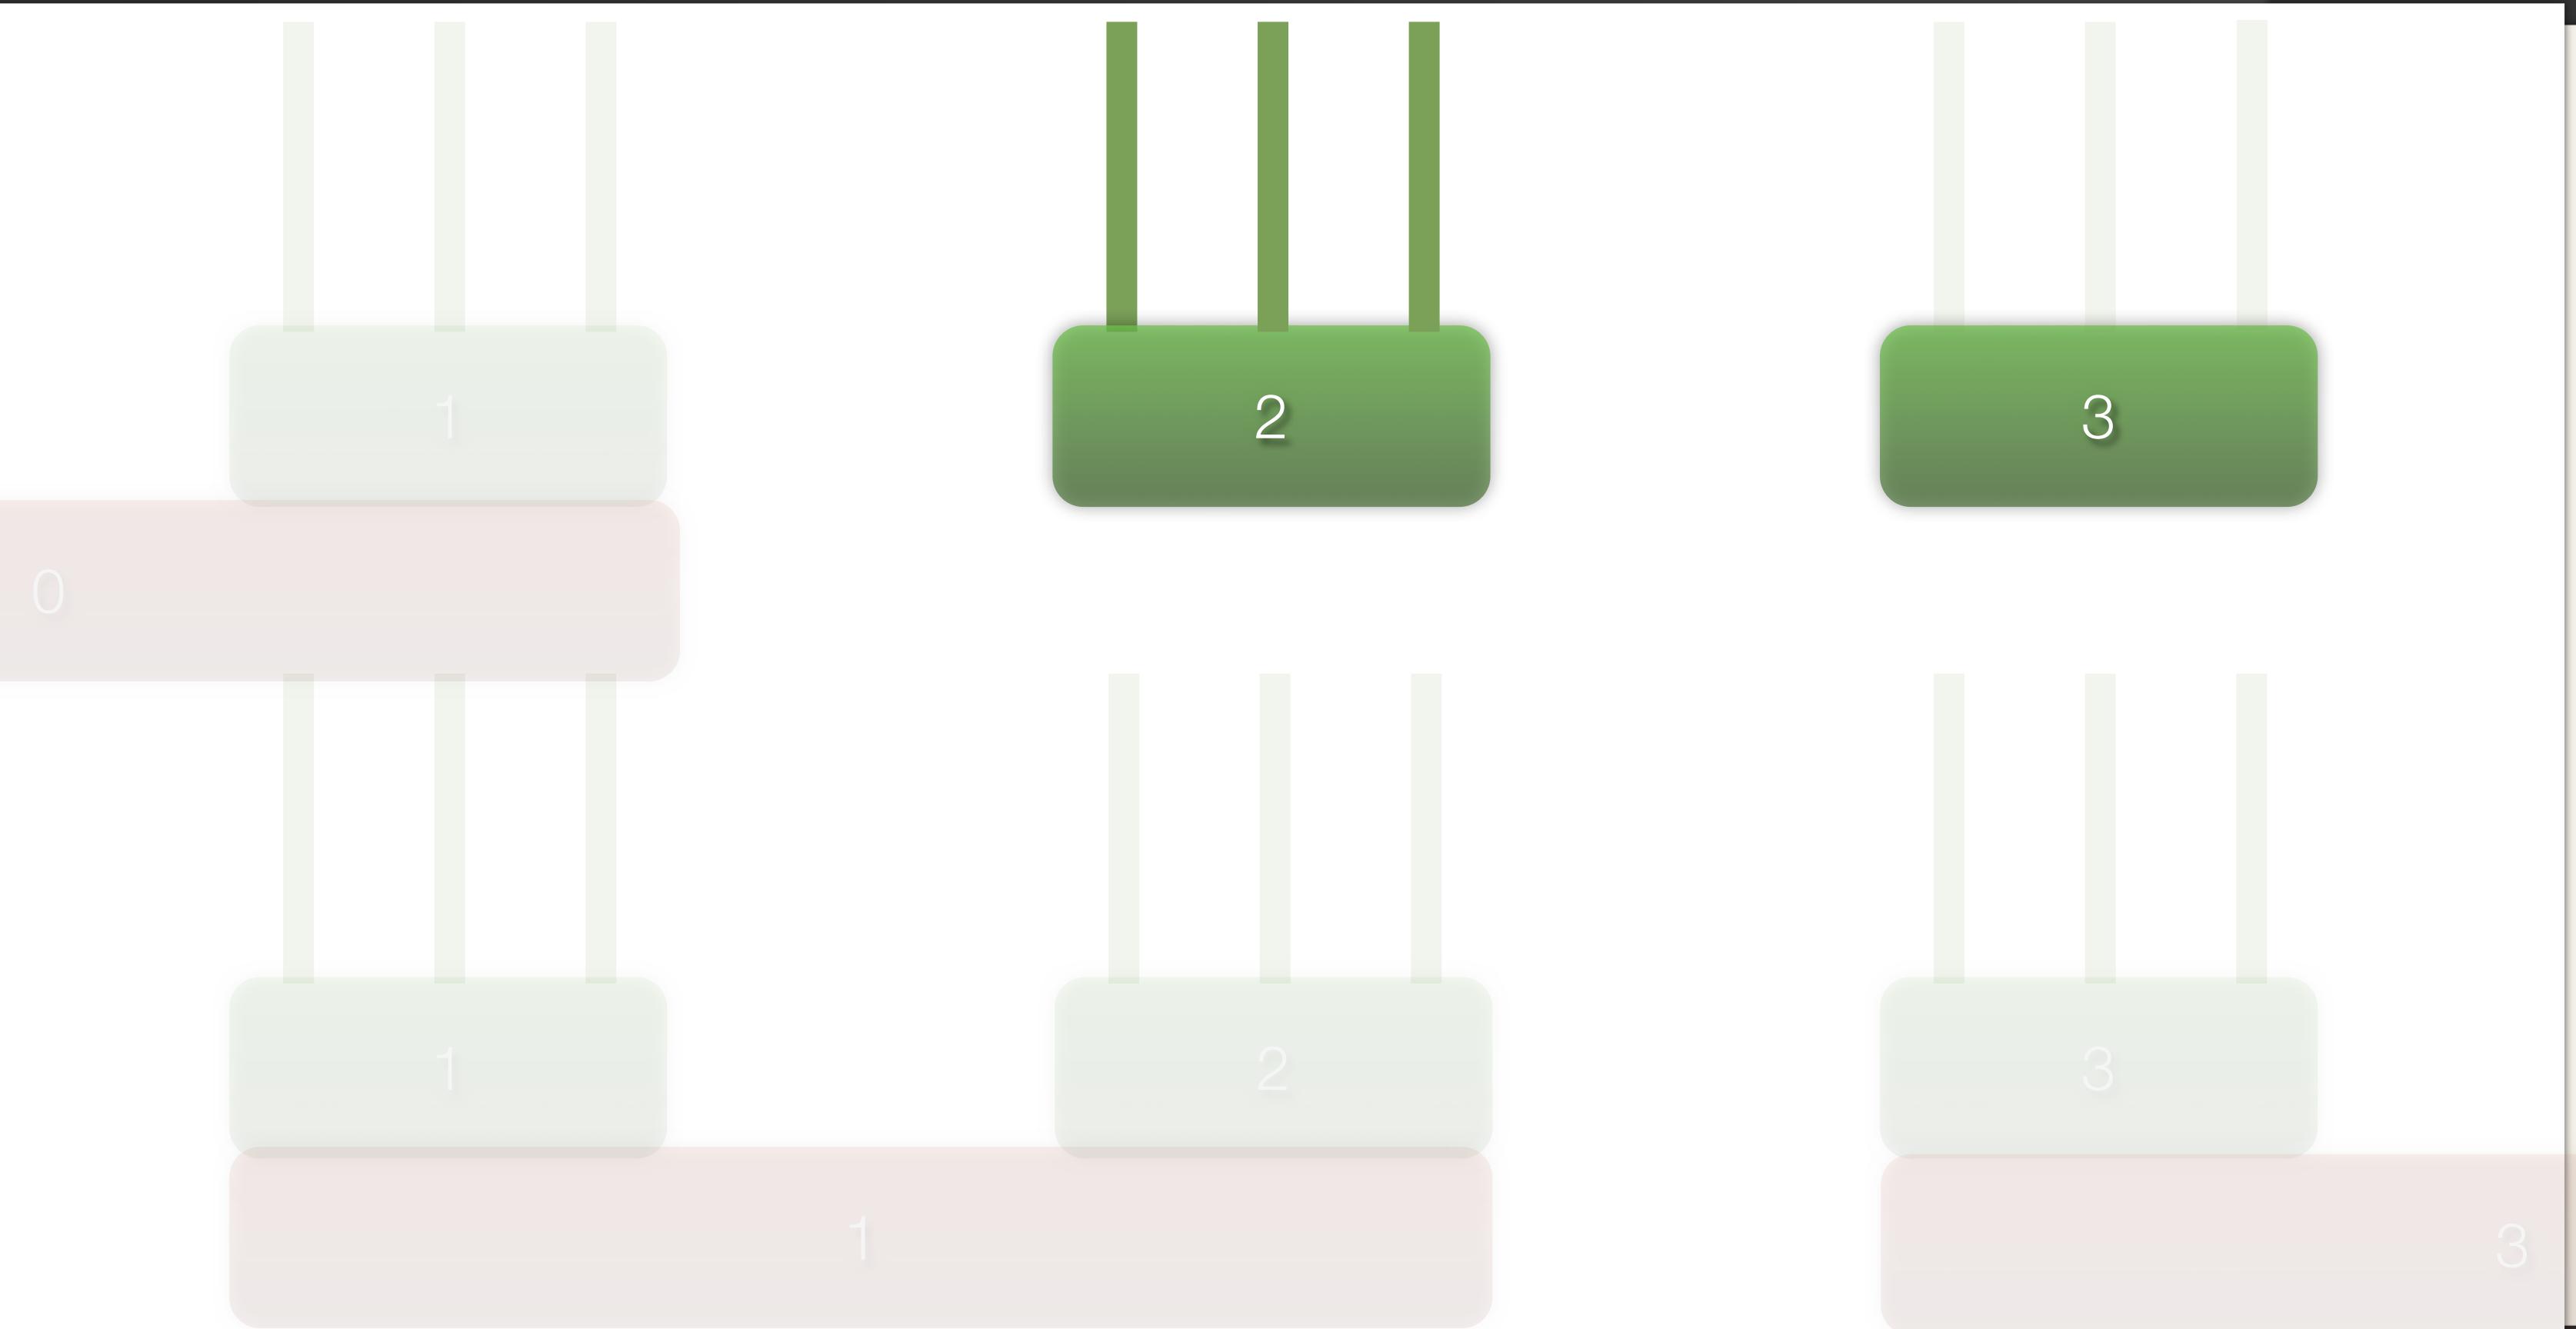


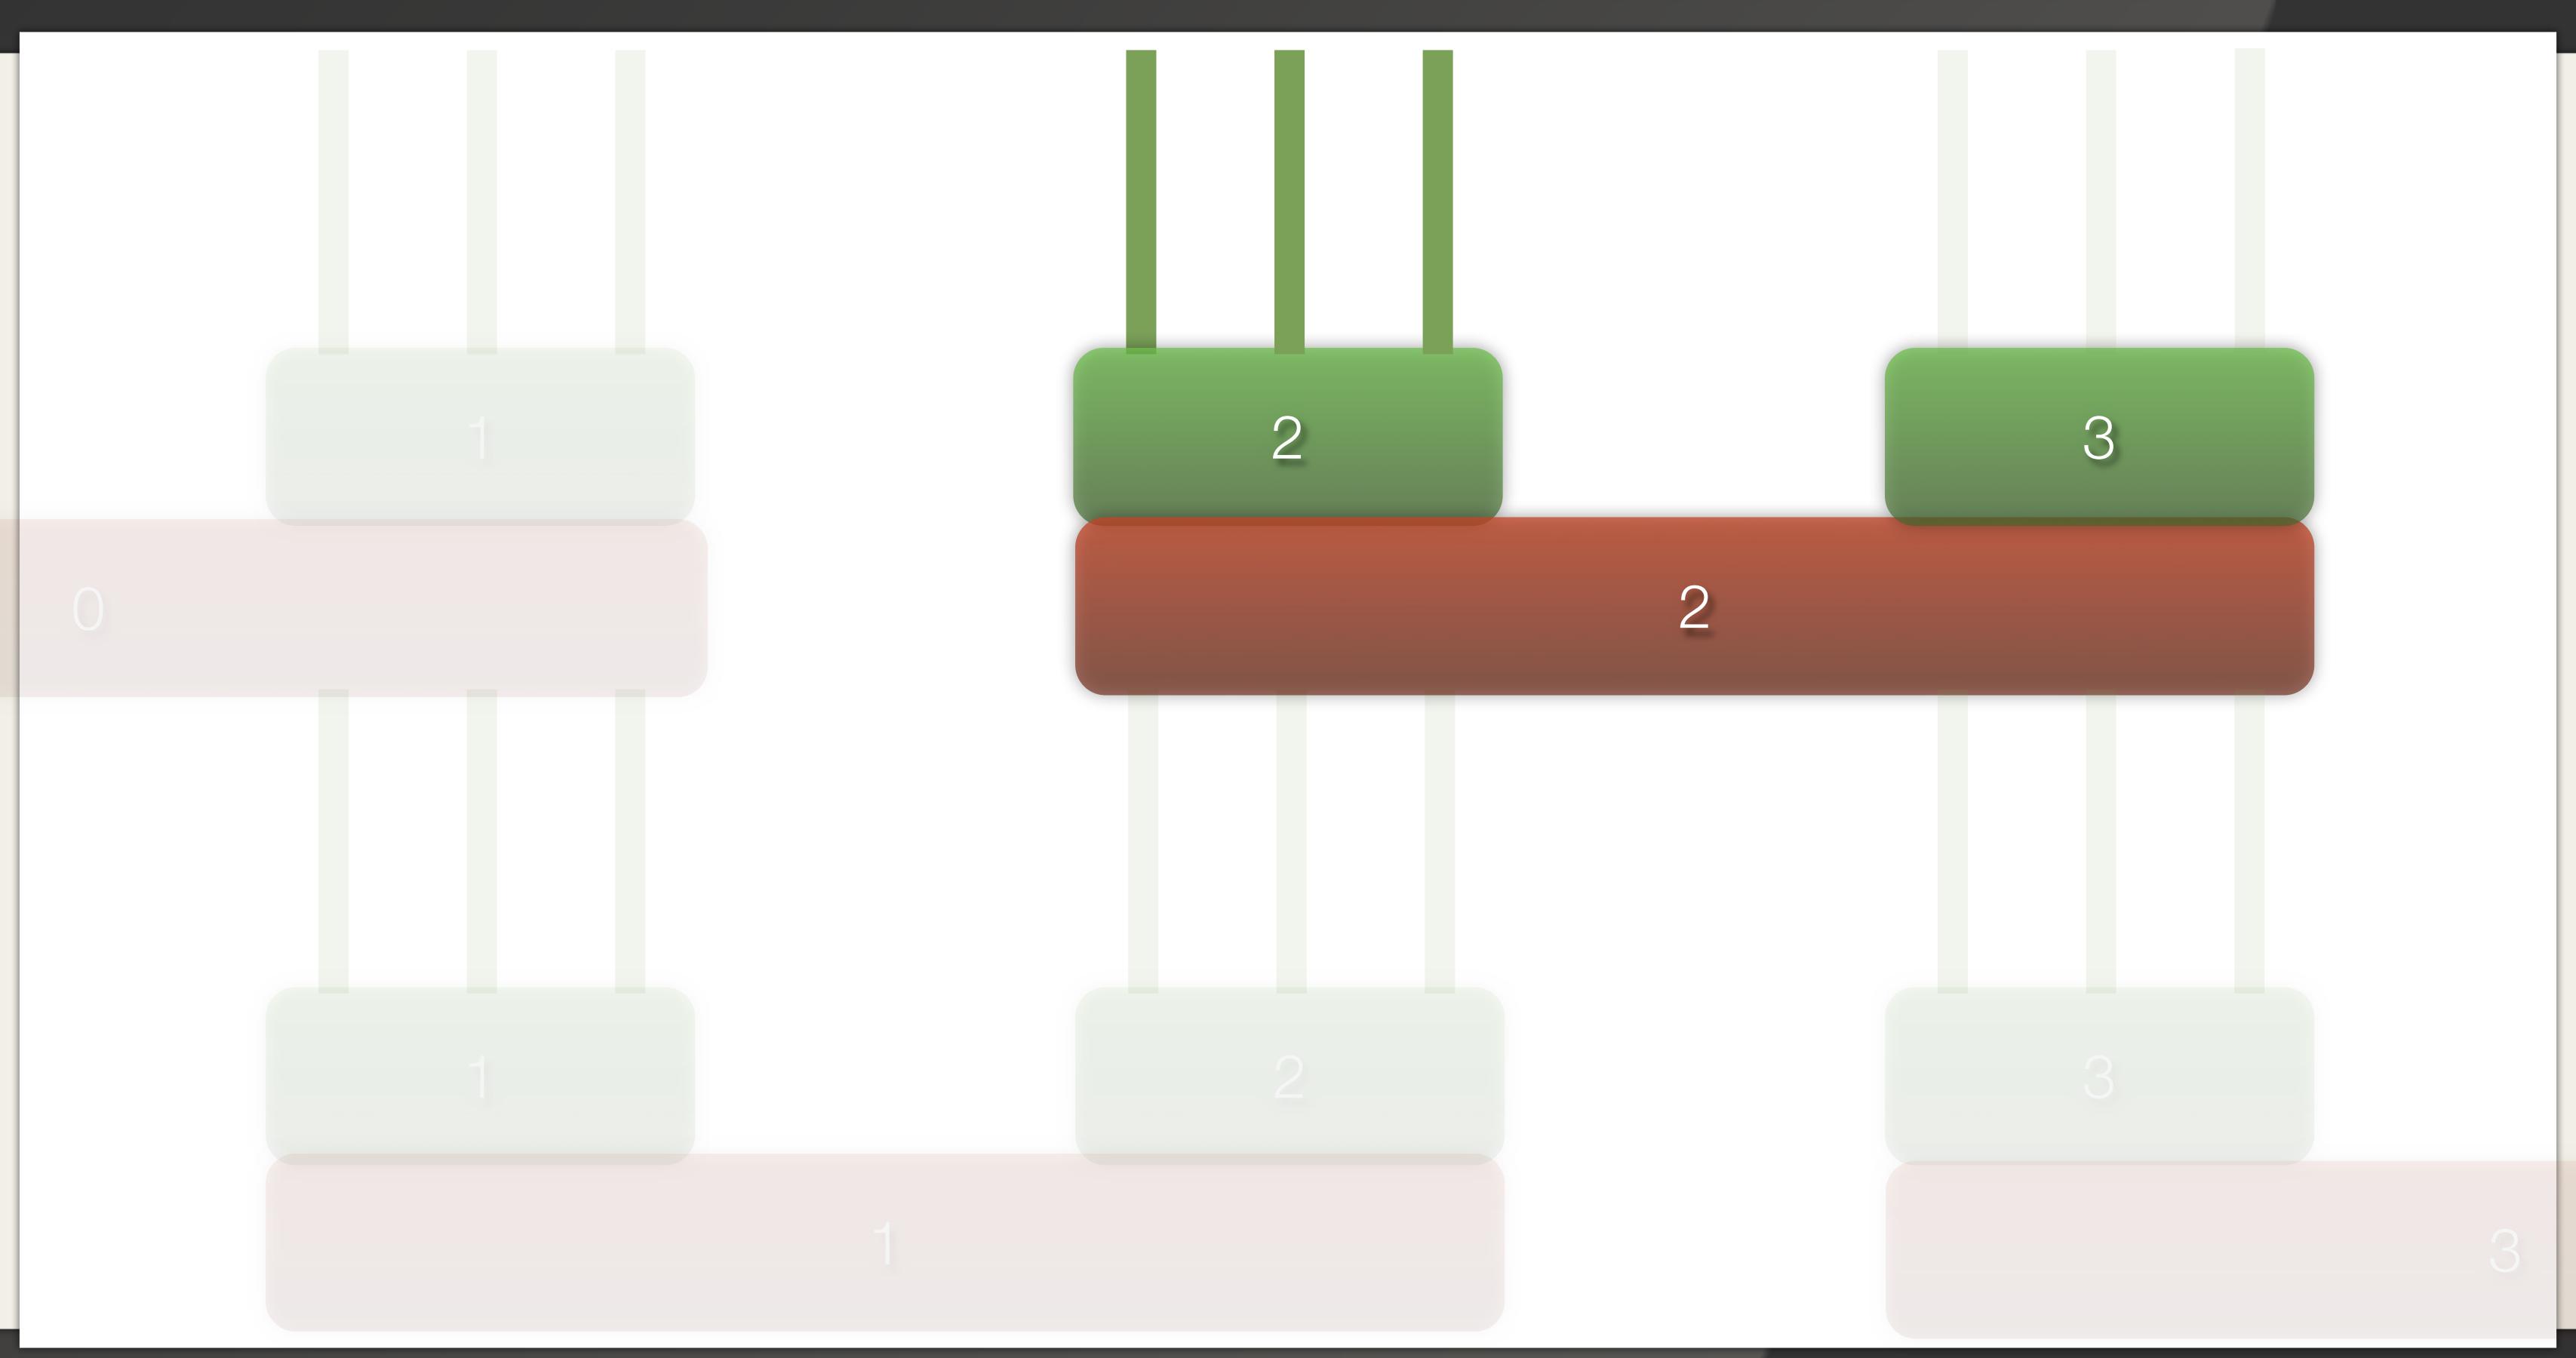


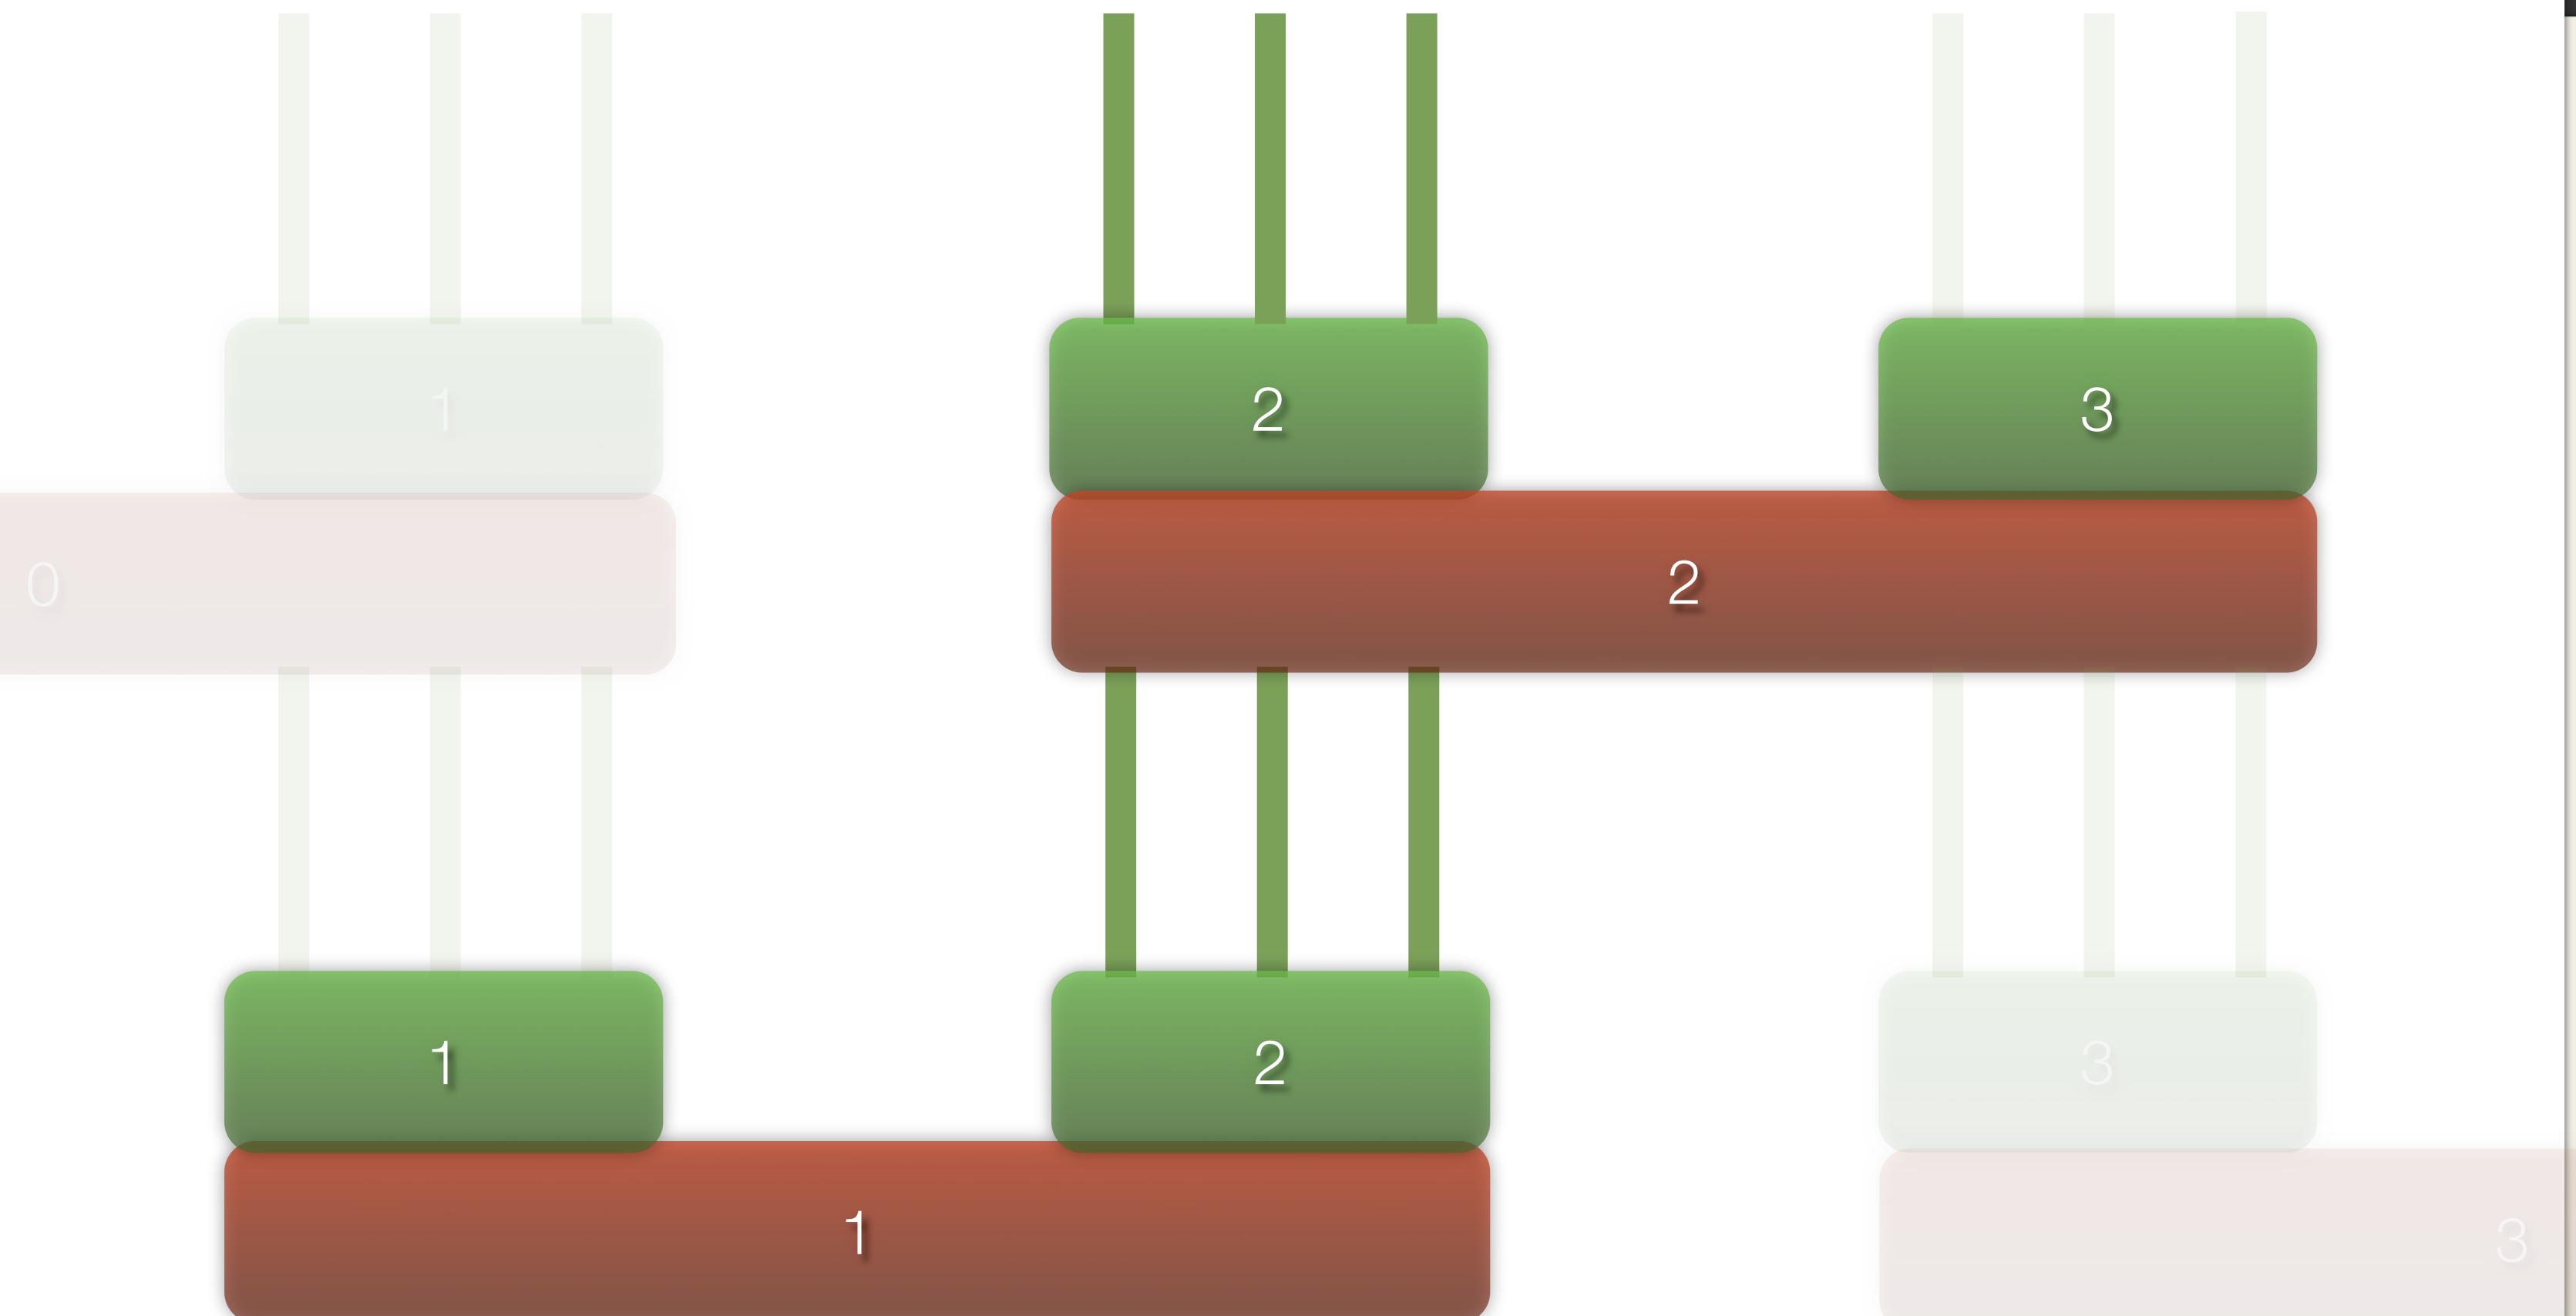


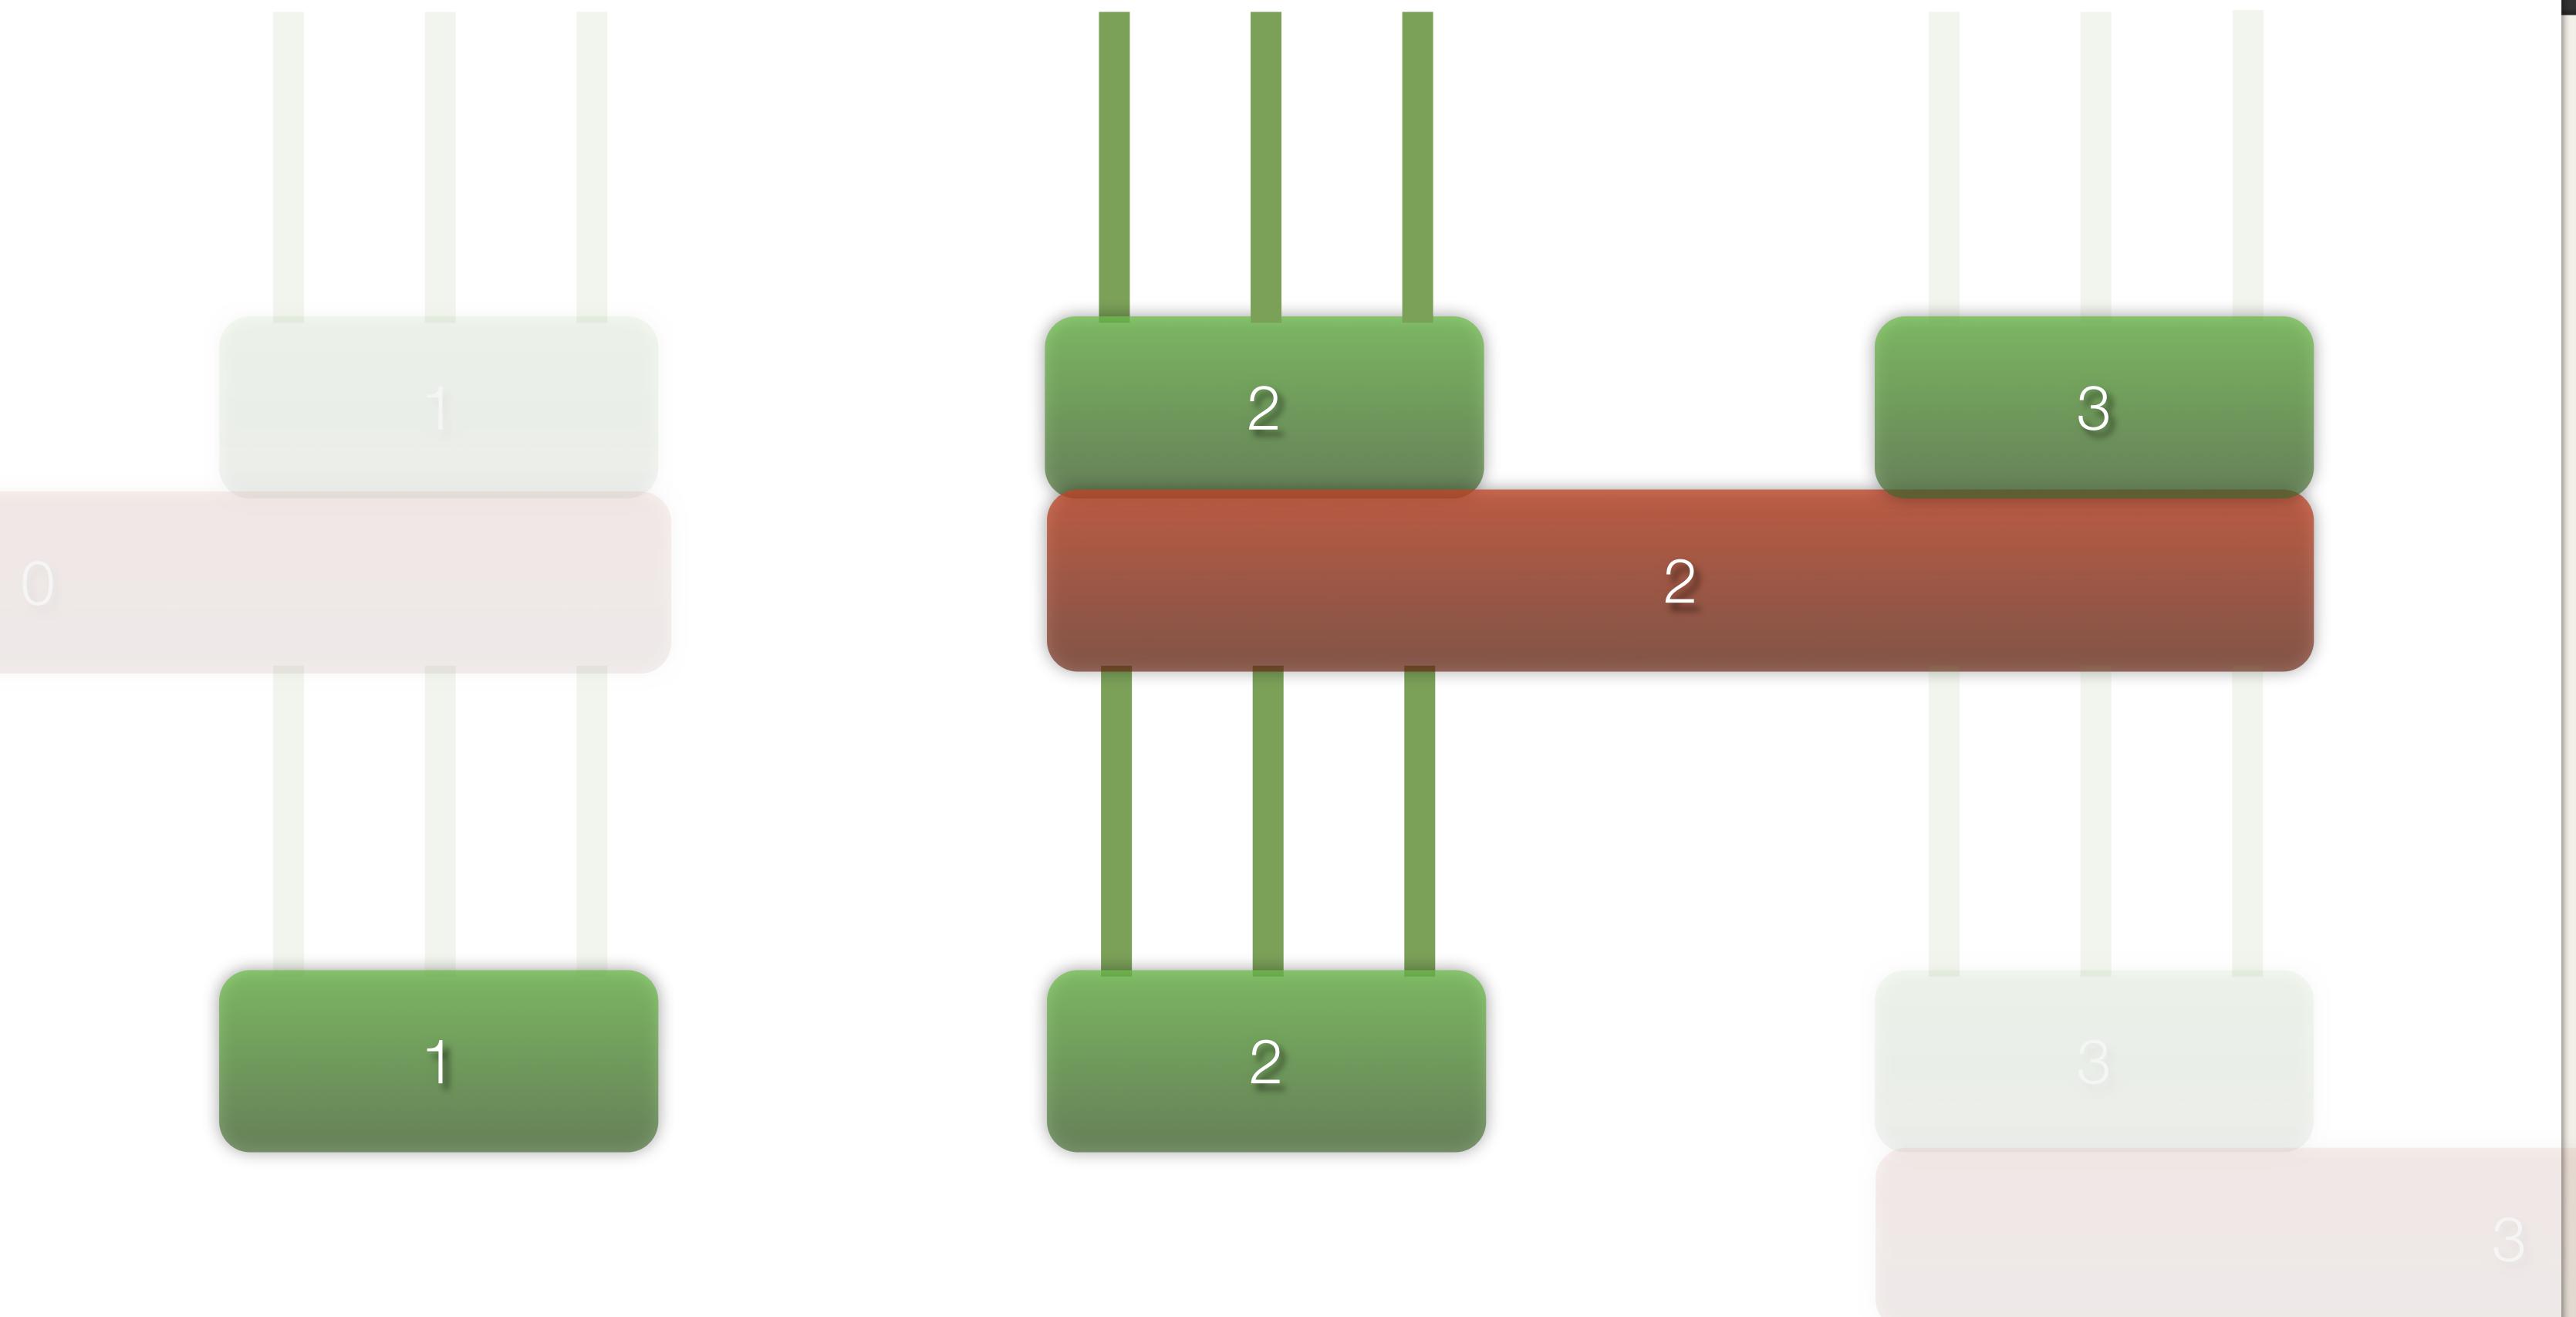












# MPI Forum Proposal

- Proposed new chapter for the MPI standard
  - Perceived as incomplete due to short timeframe of submission
  - Proposal will continue in MPI-<next>
- Research is ongoing in multiple implementations

